

Radiance Fields: Some Novel Views



THE UNIVERSITY OF
SYDNEY

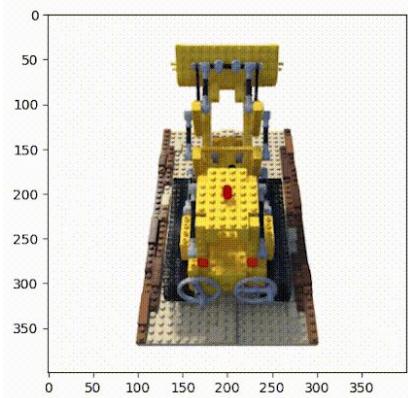
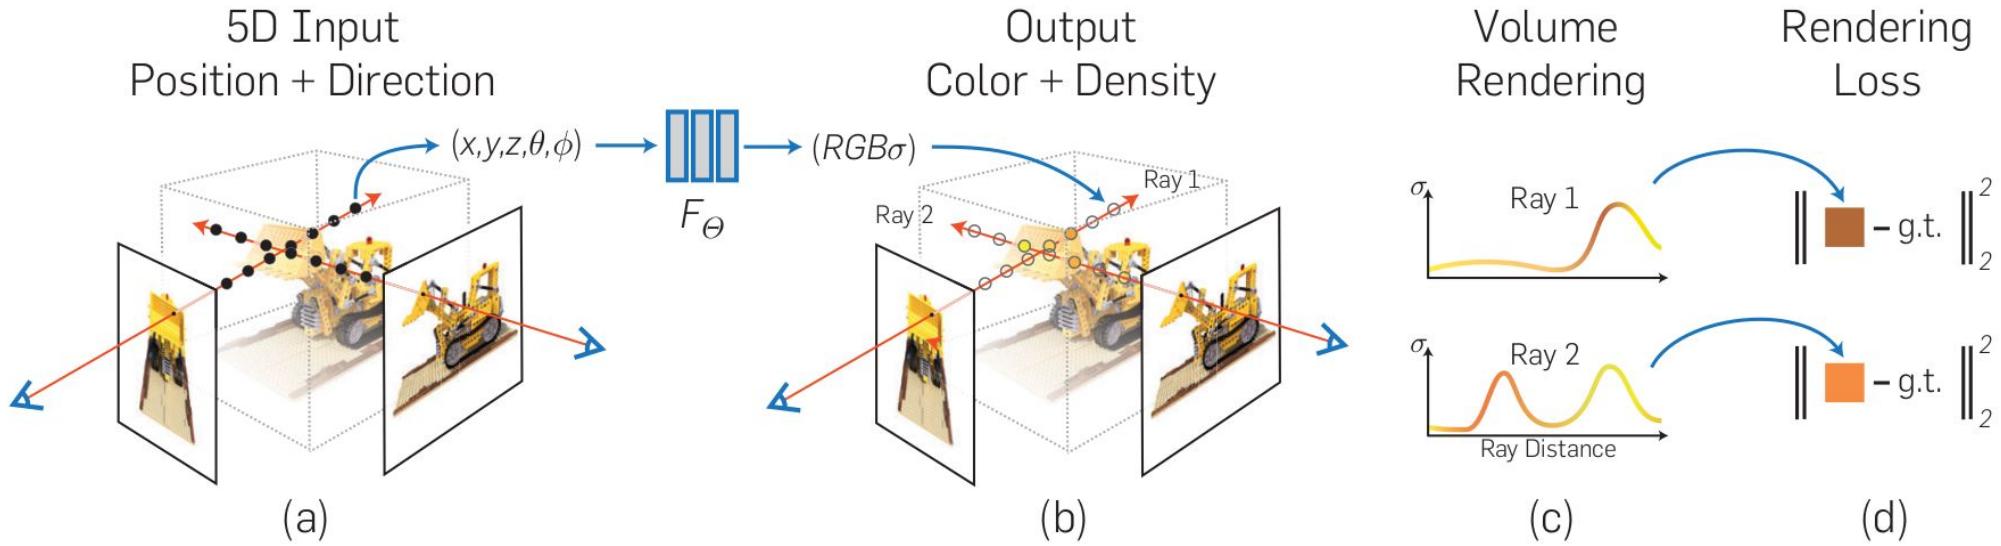


ROBOTIC
IMAGING
LAB

Don Dansereau
RVSS 2025
2025.Feb.05

Neural Radiance Fields: NeRFs

[Mildenhall2020]

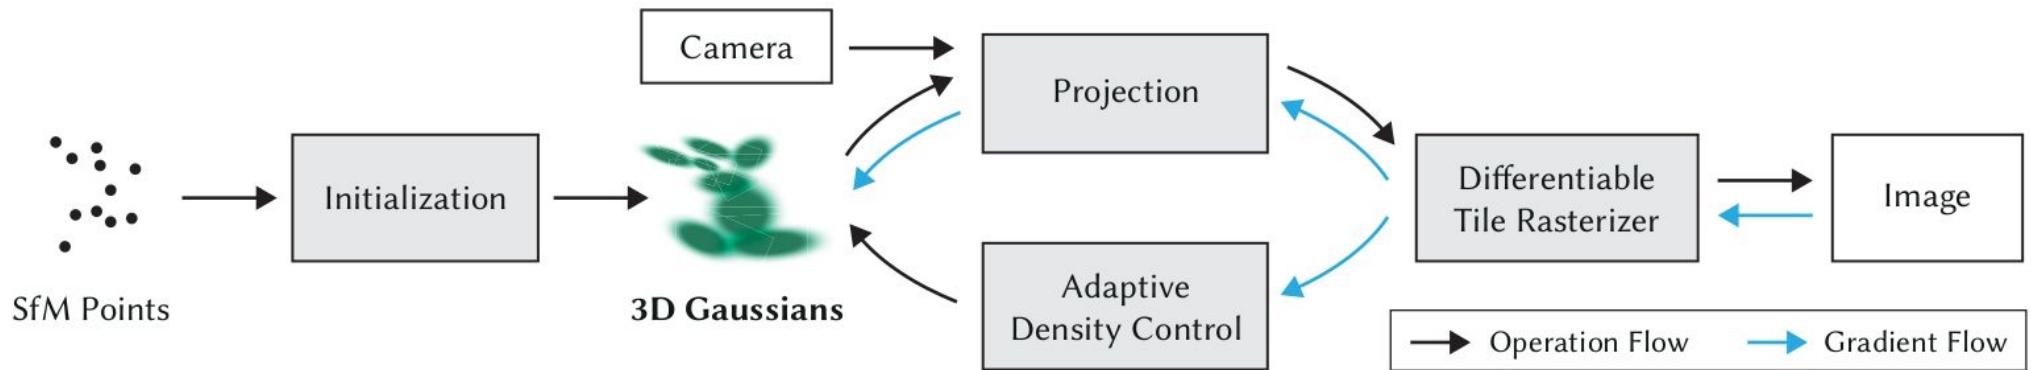


In: collection of images
Out: render novel views
Includes complex appearance
Transparency, reflection, etc.



Gaussian Splatting: 3DGS

[Kerbl2023]



In: collection of images
Out: render novel views
Includes complex appearance
Transparency, reflection, etc



Exciting Possibilities for Robotics

- Summarise many images compactly
- Natively handles complex visual appearance
- Photo-realistic

<https://gsplat.tech/>

This Lecture

Let's invent NeRFs and 3DGS...

Intuition for how they work

Similarities, differences, quirks, limitations

About the code examples*:

I want to run 100's of fast experiments

To build my intuition through tweaking and interaction

On a 5 year old laptop, no GPU, limited internet

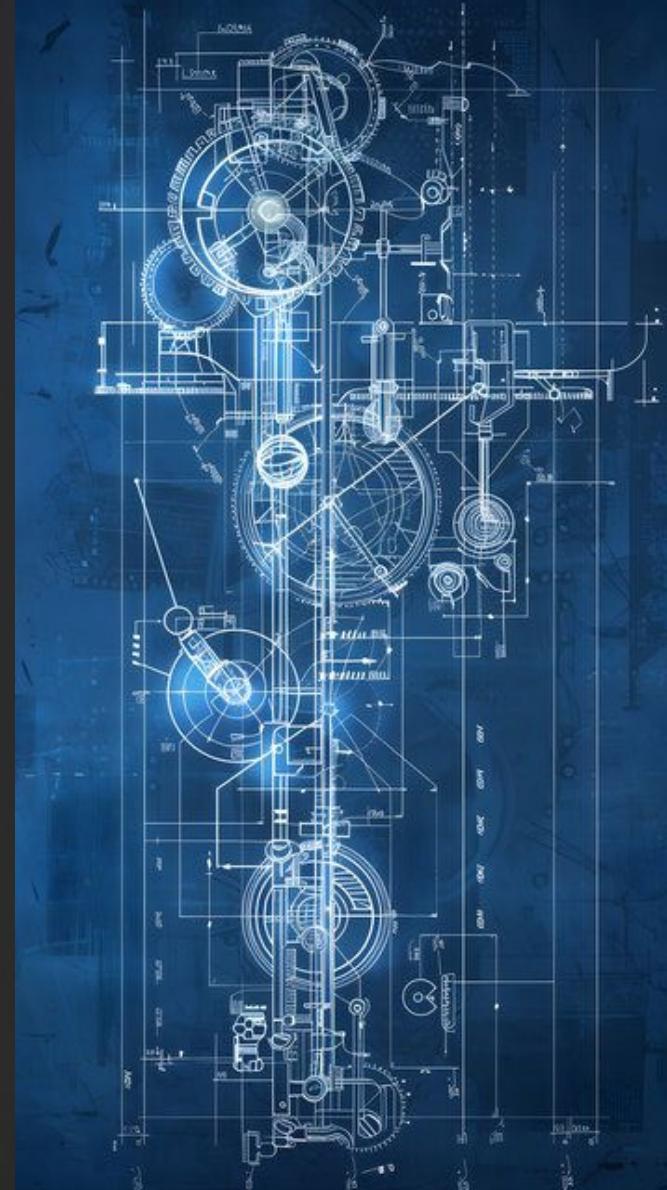
How?

Low-res, 2D, incomplete, ugly, toy examples

Not pretty, not state of the art...

... but that's not the point.

*with credit /blame to ChatGPT



From the Top

What kind of algorithm is this?

Given a set of input images

Adjust a model

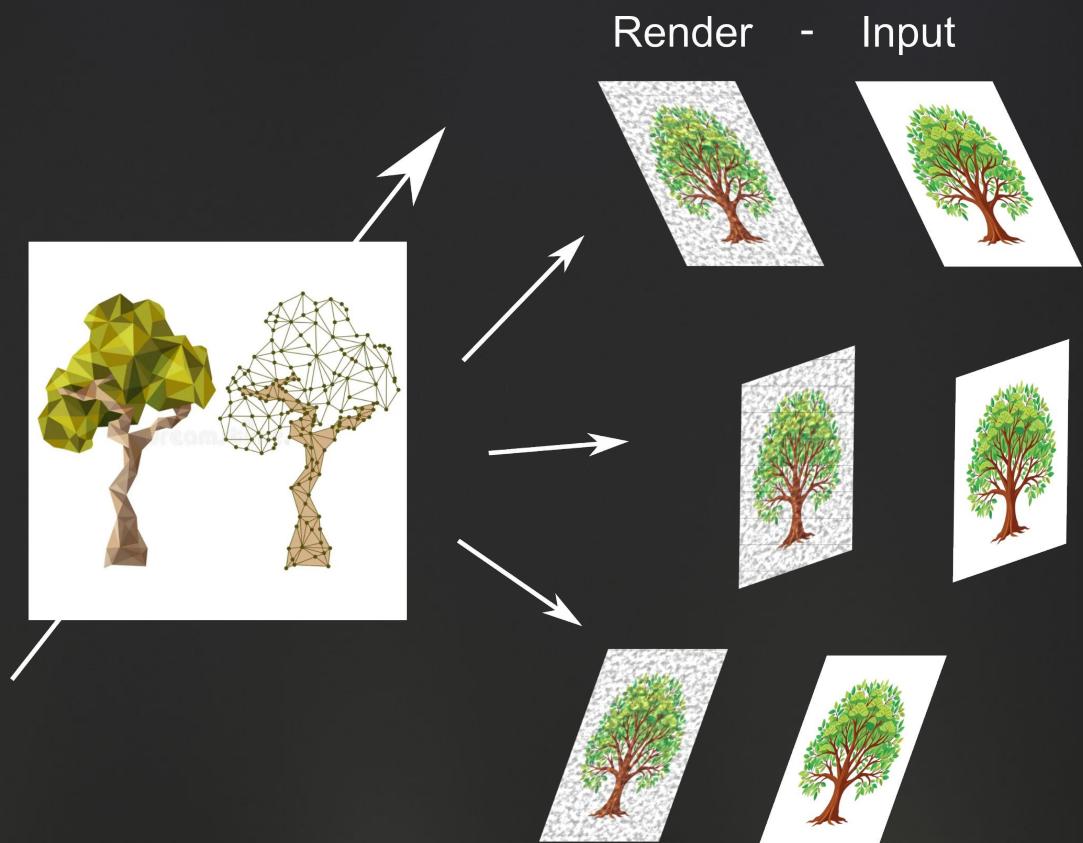
Until renders match inputs

“Inverse rendering”

e.g. Mitsuba

Typically polygonal, “hard” models

Often the model is the point



From the Top

What kind of algorithm is this?

Given a set of input images

Render novel views

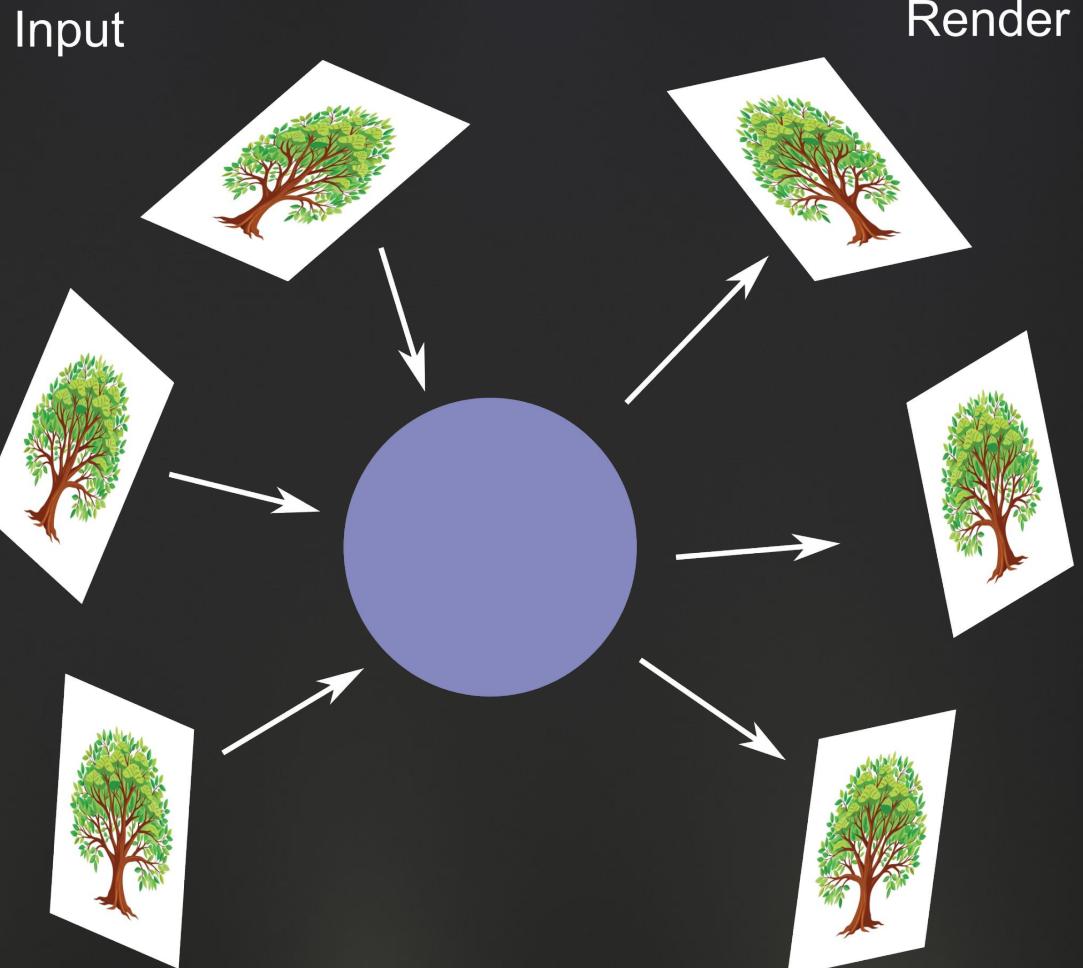
Directly from the images

“Image-based rendering”

e.g. Bullet time effect, The Matrix

Usually novel views are the point

“Novel view synthesis”



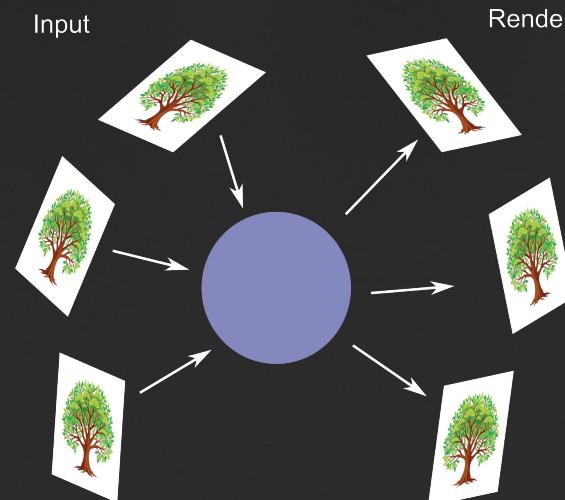
NeRF and 3DGS?

Given a set of input images

Render novel views

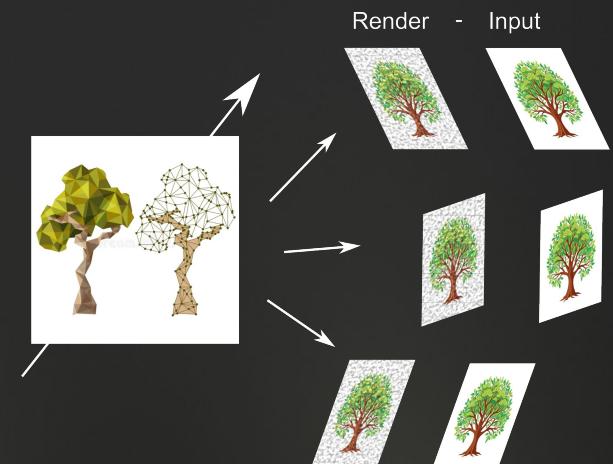
By fitting a sort of a model...

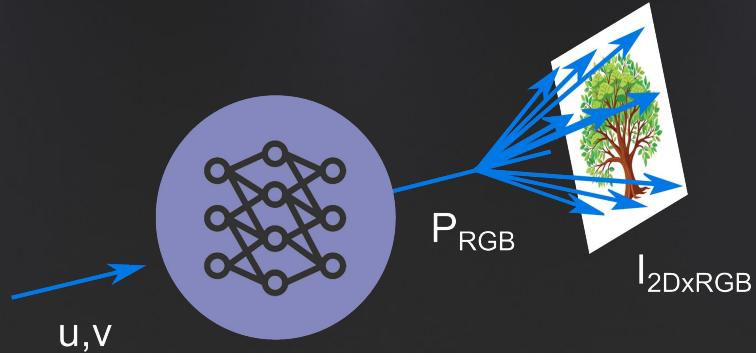
... but not a hard model



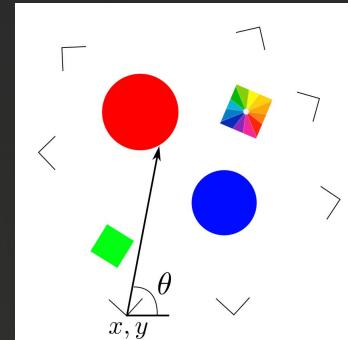
Both image-based and inverse rendering

... but with something special inside.





P1: Inventing Inverse Rendering in 2D



P2: Inventing Novel View Synthesis in Flatland

Inventing Neural Inverse Rendering for 2D Images

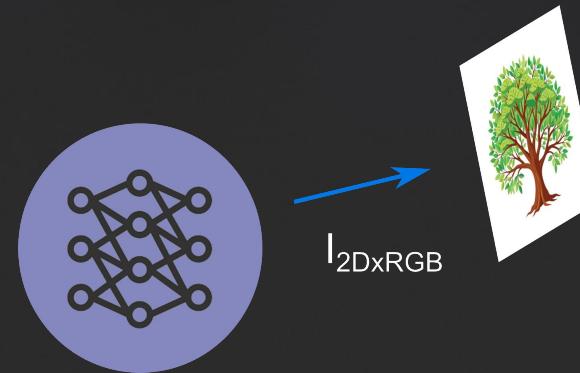
Approximate an RGB image $I(u,v)$

NN with big output layer

Directly builds 2D array of RGB pixels

Hard to learn

Not that useful (e.g. interpolate?)



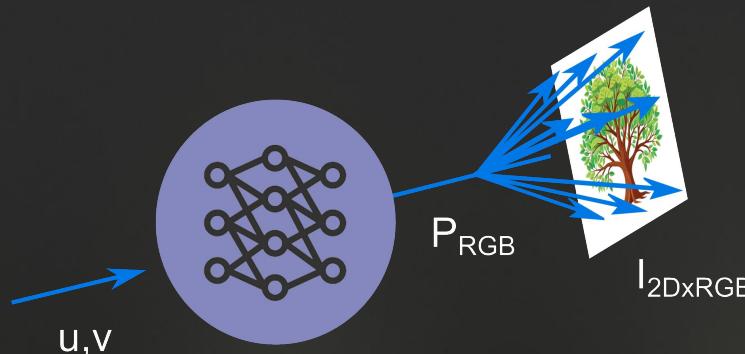
NN with small output layer

Single RGB pixel output

We provide the coordinate

We do more of the work

More useful (e.g. interpolate)



Inventing Neural Inverse Rendering for 2D Images

`invrender_2d_nn_simple.py`

It kind of works!

Struggling with fine detail

Our input is very smooth

v,u:

[0.0, 0.0], [0.0, 0.1], [0.0, 0.2], ...

[0.1, 0.0], [0.1, 0.1], [0.1, 0.2], ...

[0.2, 0.0], [0.2, 0.1], [0.2, 0.2], ...

...

More informative input?



Output



Target

Inventing Positional Encoding

Let's provide more texture in the input
Helps the network distinguish neighbours

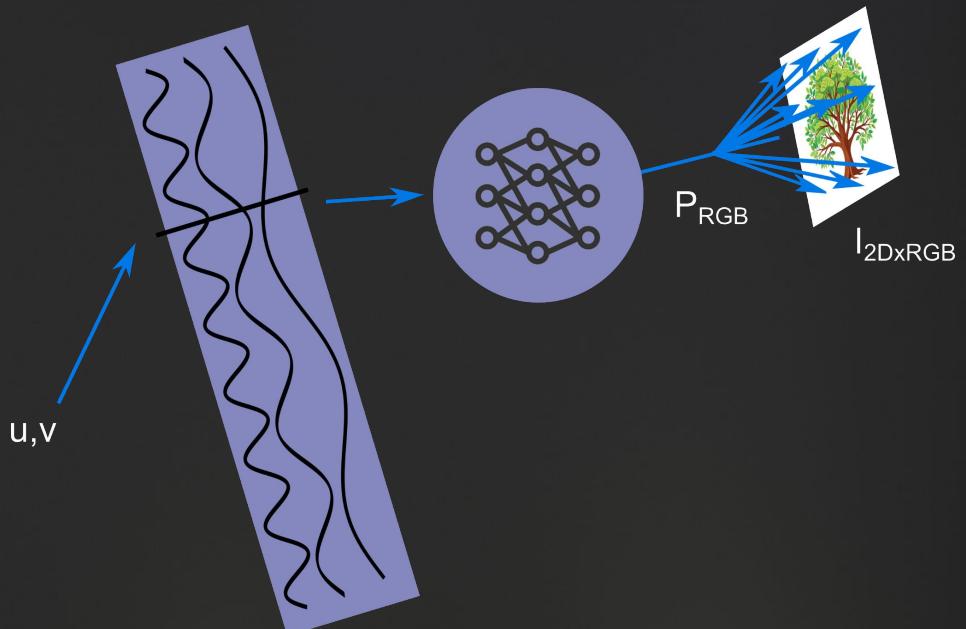
u, v

$\cos(u), \cos(2u), \cos(4u), \dots$

$\sin(u), \sin(2u), \sin(4u), \dots$

$\cos(v), \cos(2v), \cos(4v), \dots$

$\sin(v), \sin(2v), \sin(4v), \dots$



Neighbours are distinct, unique
(*smells a little Fourier?*)

Inventing Neural Inverse Rendering for 2D Images

invrender_2d_nn_positional.py

Much better with positional encoding!

Try:

Twiddle every single parameter

Interpolate / upsample

Learn multiple images

Incrementally build a panorama



Output



Target

Inverse Rendering 2D Images with Point Primitives

Build image from many simple primitives

Adjust primitives to match target

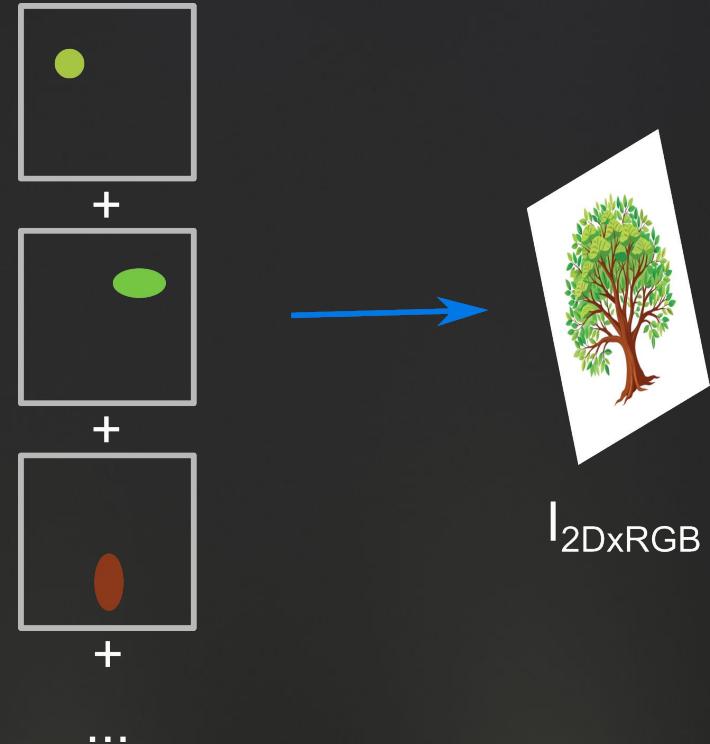
Let's start with many circles

`invrender_2d_circles.py`



Spatial gradient is not smooth

No slope to follow

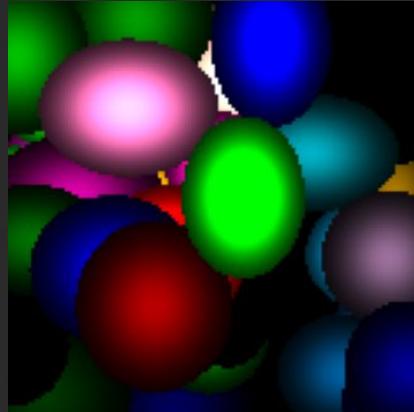


Inverse Rendering 2D Images with Gaussian Primitives

Again with opaque Gaussians

Adjust `invrender_2d_gaussians.py`

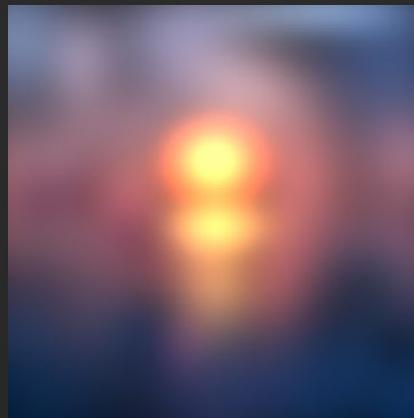
Smooth edges allow some gradient
... but most changes are hidden in back



Again with transparent Gaussians

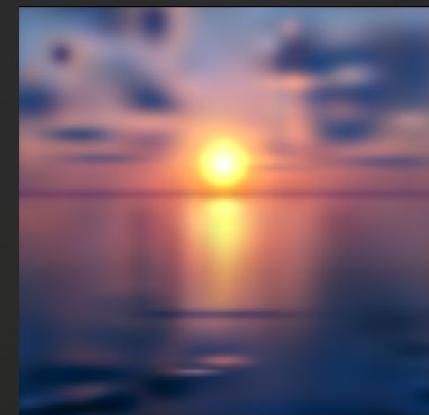
Adjust `invrender_2d_gaussians.py`

Adding up all Gaussians at each pixel



Much more promising!

Again with 128 blobs



Recap: Inverse Rendering

Many representations can capture the same scene

NNs benefit from input coordinates and positional encoding

Inverse rendering needs soft representations

Provides a gradient to climb

Watching the optimisation process can build intuition

Ideas to try:

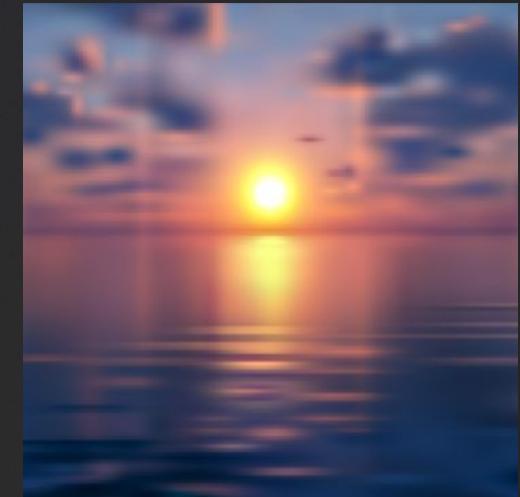
NN hash encoding (sparse and dense examples)

Gaussian blobs with rotation

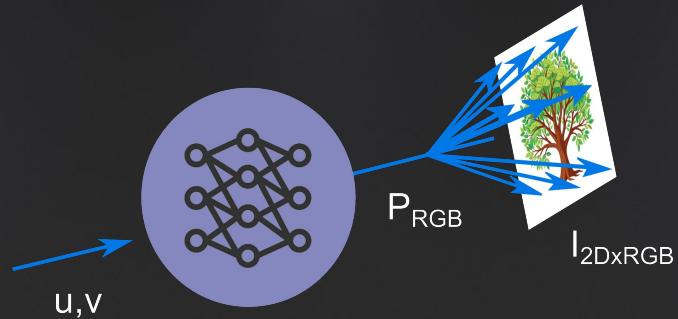
Periodically add / remove blobs / reduce opacity of blobs

Regularisers!

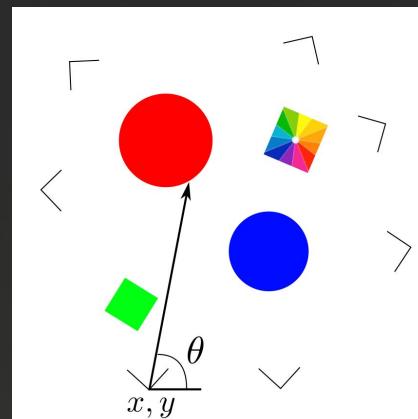
Video; video interpolation!



Gaussians
w/dynamic addition
/ removal



P1: Inventing Inverse Rendering in 2D



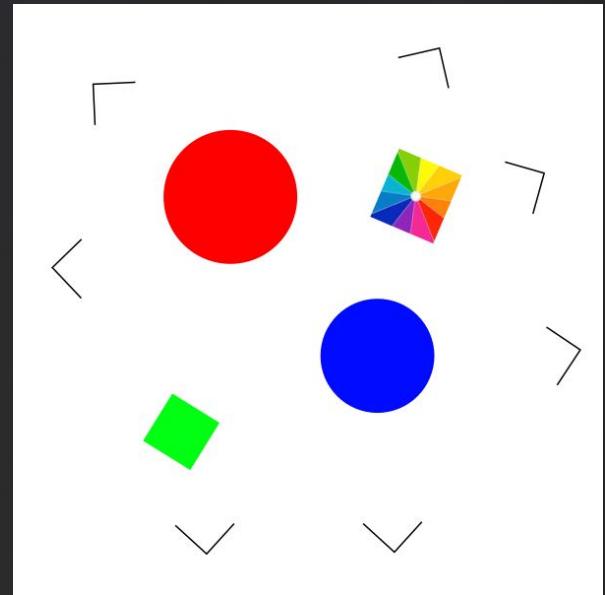
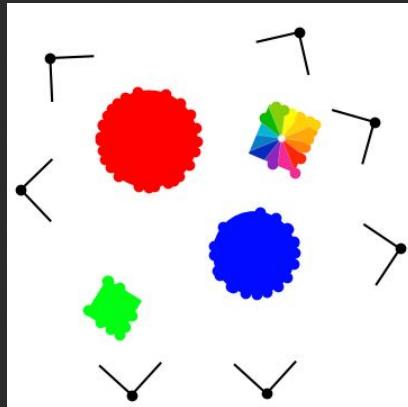
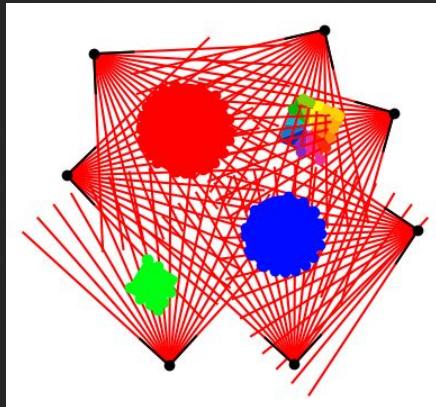
P2: Inventing Novel View Synthesis in Flatland

Welcome to Flatland

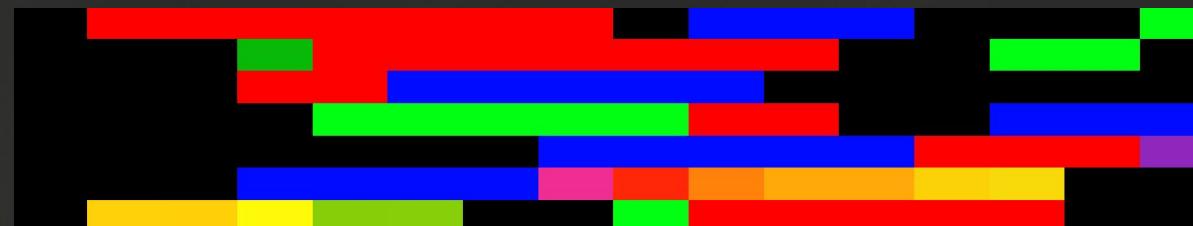
What does a camera see in flatland?

`renderer2d.py`, `flatland_intro.svg`

Loads and renders from SVG files



Cameras



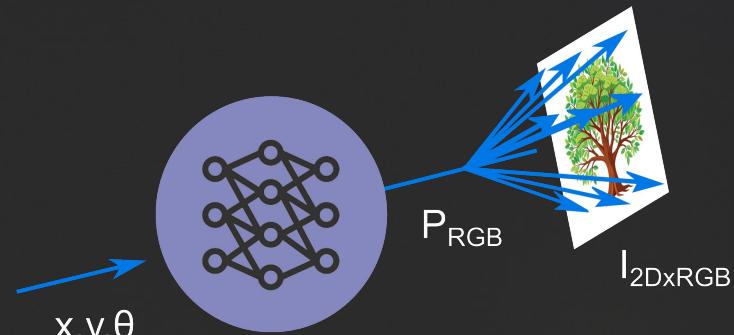
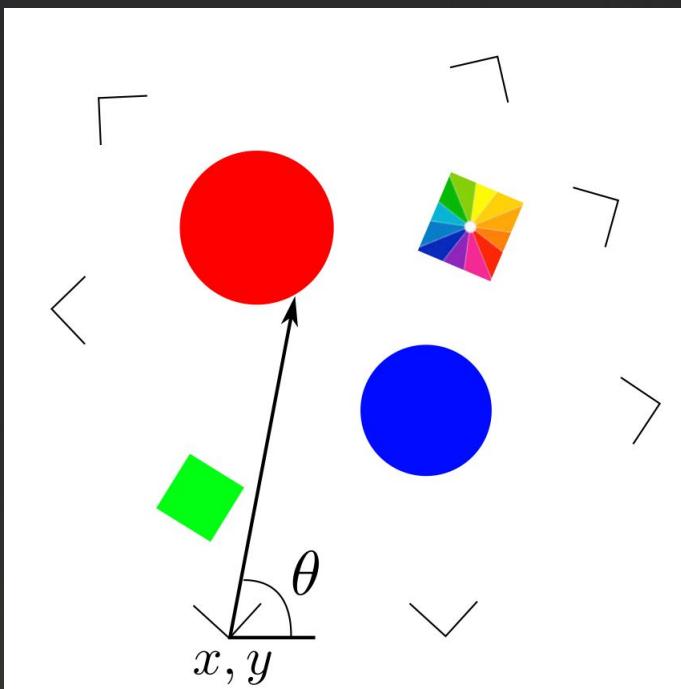
Pixels

Inverse Rendering from Rays

Let's do inverse rendering on the space of light rays

Like our 2D renderer on the space of pixels

How to describe each light ray?

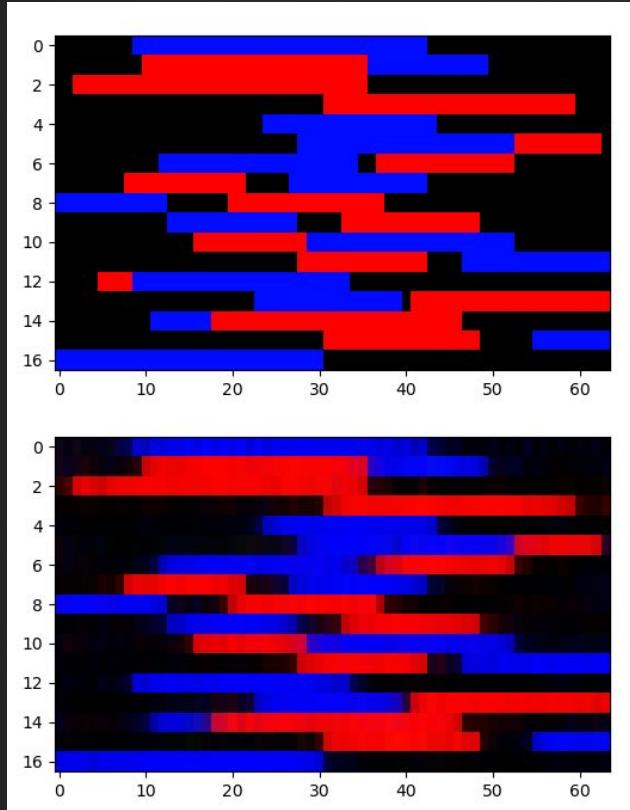


+Positional encoding

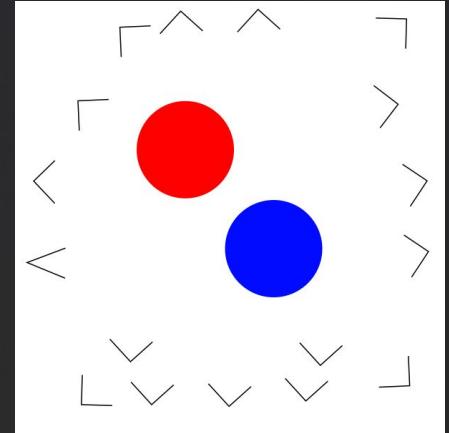
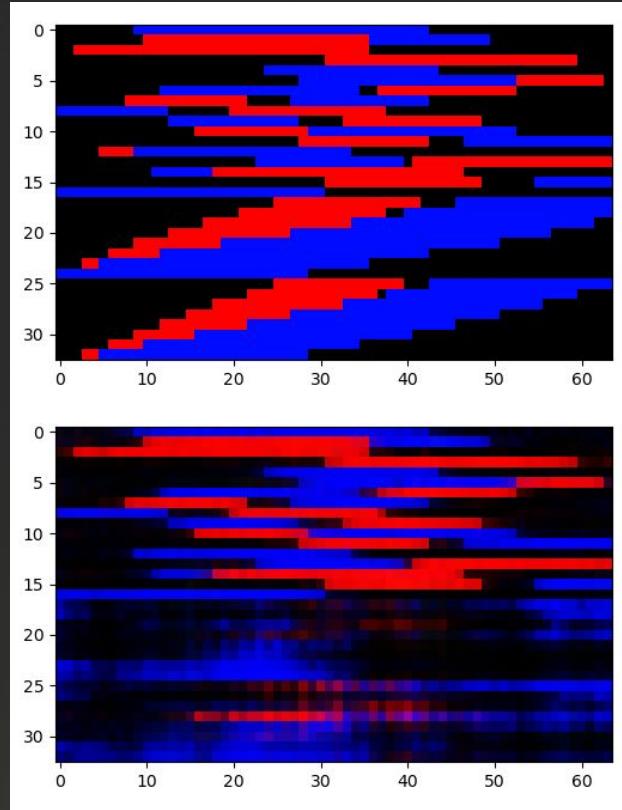
Inverse Rendering from Rays

nvs_nn_full_ray.py, scene_simple.svg

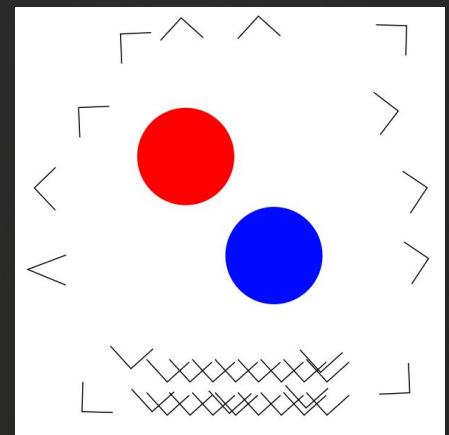
Trains fine



Test not so good



Train



Test

Let's Invent a Simpler Ray

Some rays are identical

But have different parameterizations in x, y, θ

New ray representation:

Same identity along direction of propagation

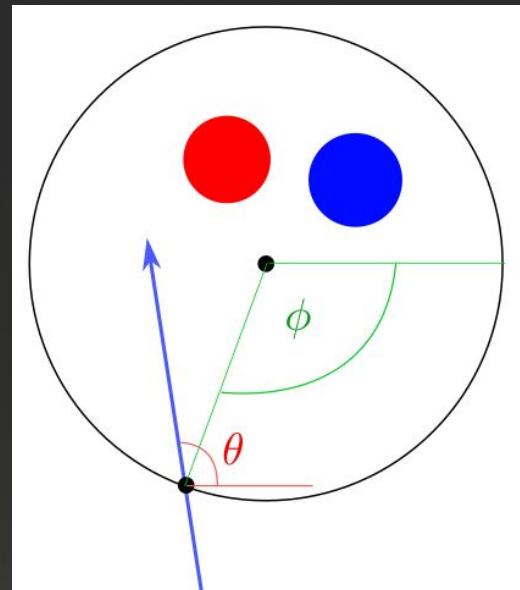
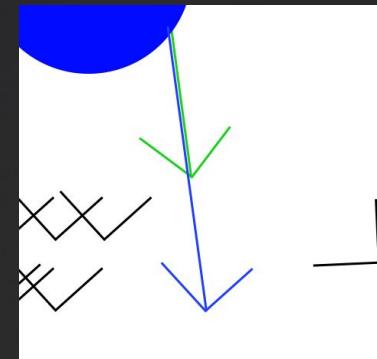
Reference surface around scene

Point of intersection on the surface ϕ

Direction θ

Overlapping rays get same parameterisation

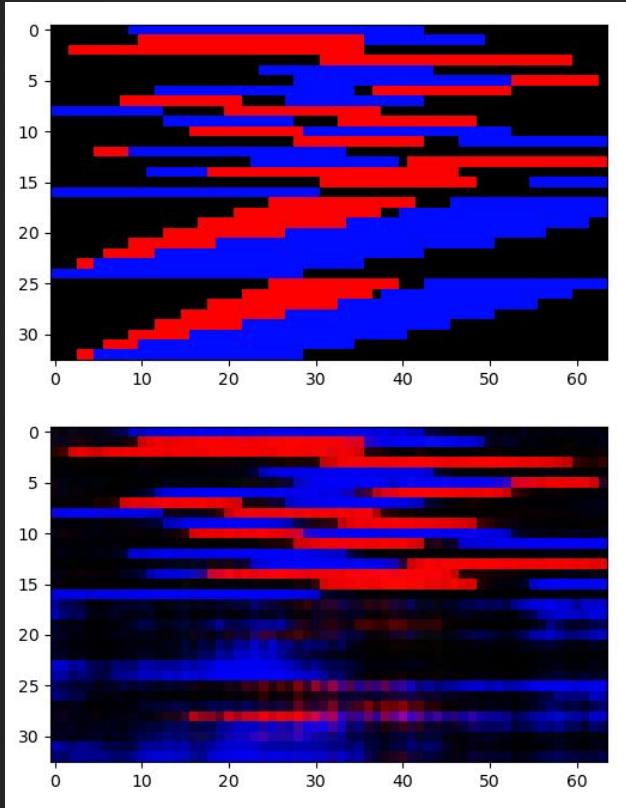
Fewer dims: learning in 2D should be easier...



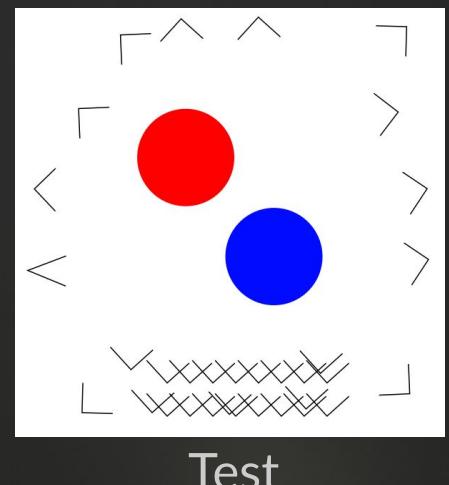
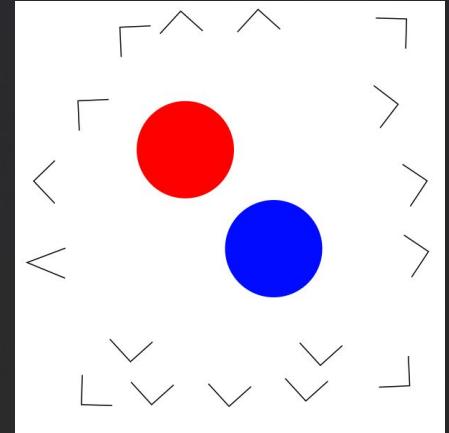
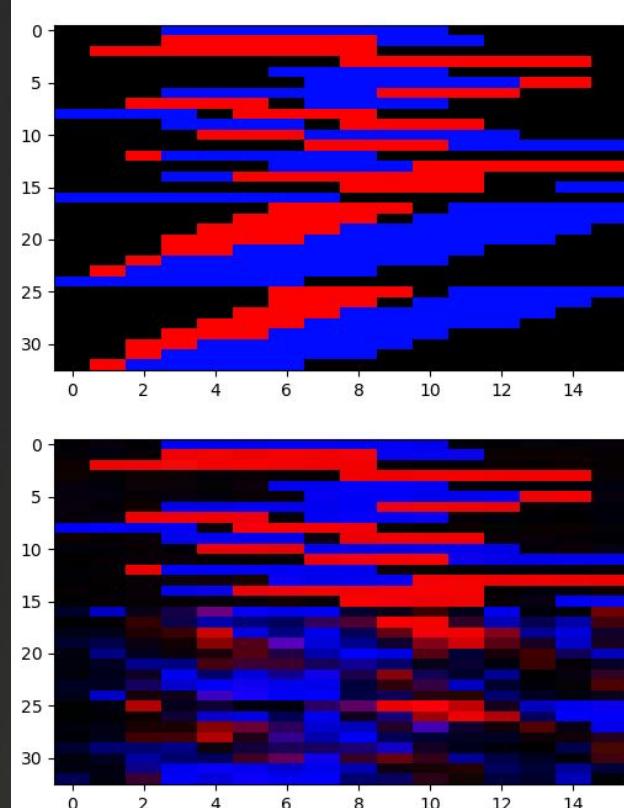
Rendering Simpler Rays

nvs_nn_lf.py

Full ray

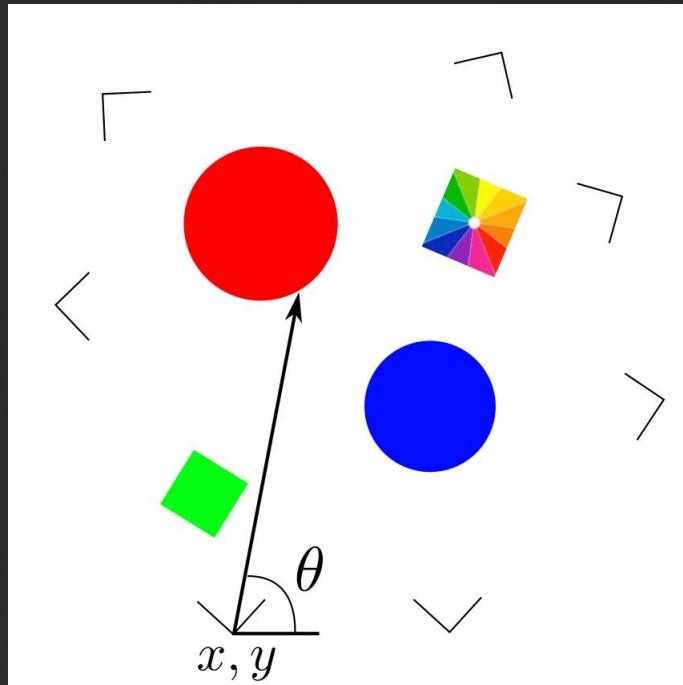


Simplified ray, not great...

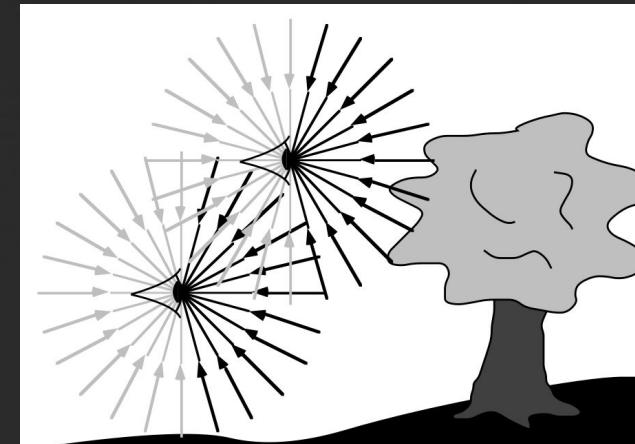


Recap and Terminology so Far

In 2D: $L(x, y, \theta)$



In 3D: $L(x, y, z, \theta, \phi)$



Adelson and Bergen, 1991
“The Plenoptic Function and the
Elements of Early Vision”

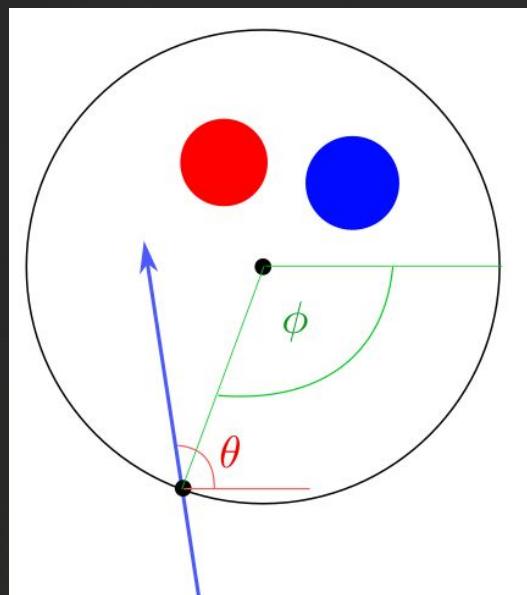
Recap and Terminology

In 2D: $L(\theta, \phi)$

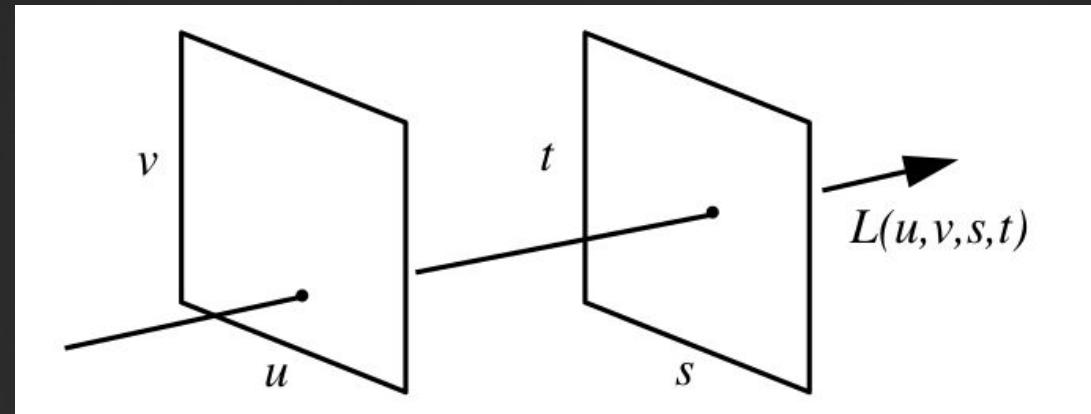
Reference surface

1D position

1D direction



In 3D: $L(s, t, u, v)$



Levoy and Hanrahan, 1996

“Light Field Rendering”

2D position

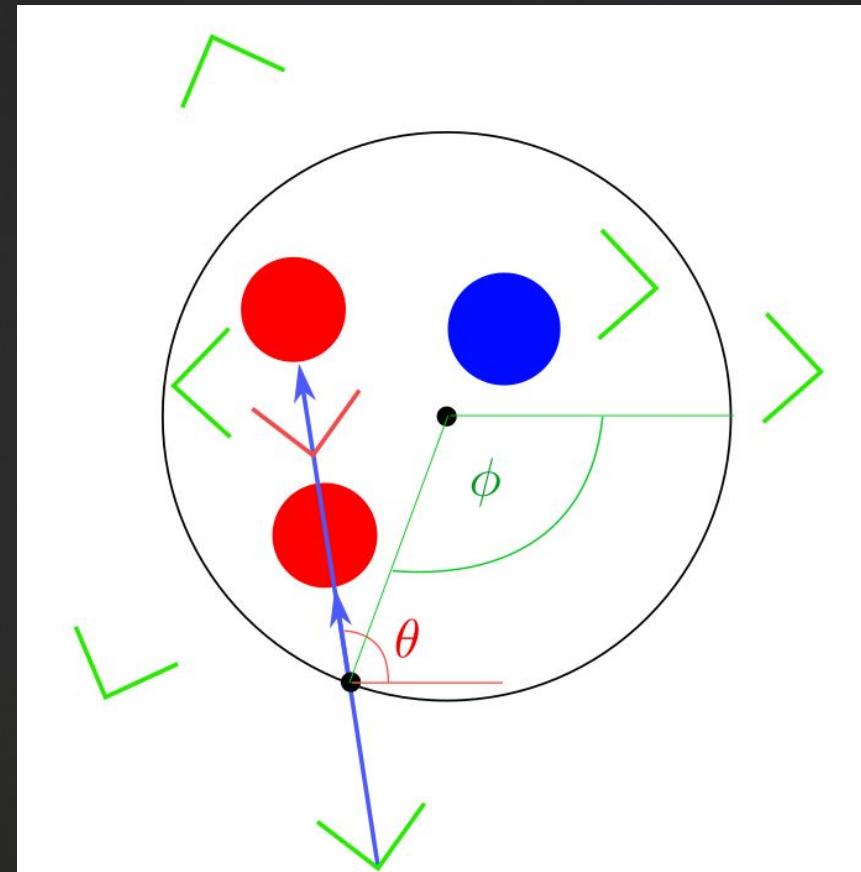
2D direction

Problems with the Light Field

Occlusions are mostly fine...

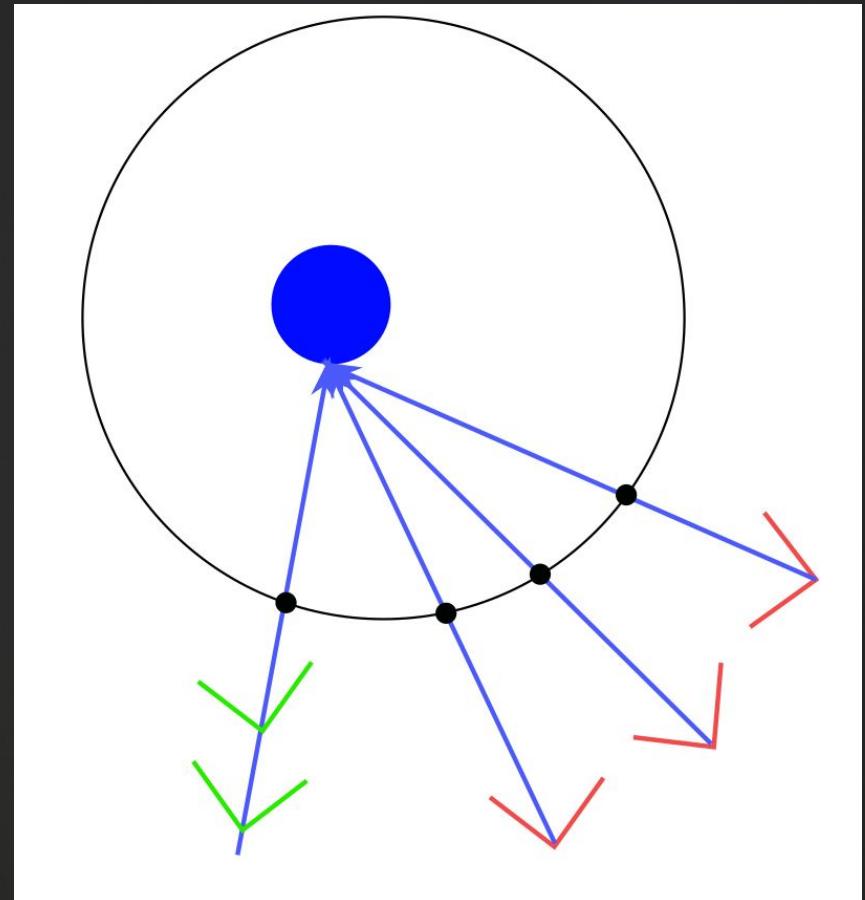
With some exceptions (e.g. in red):

One ray, two values

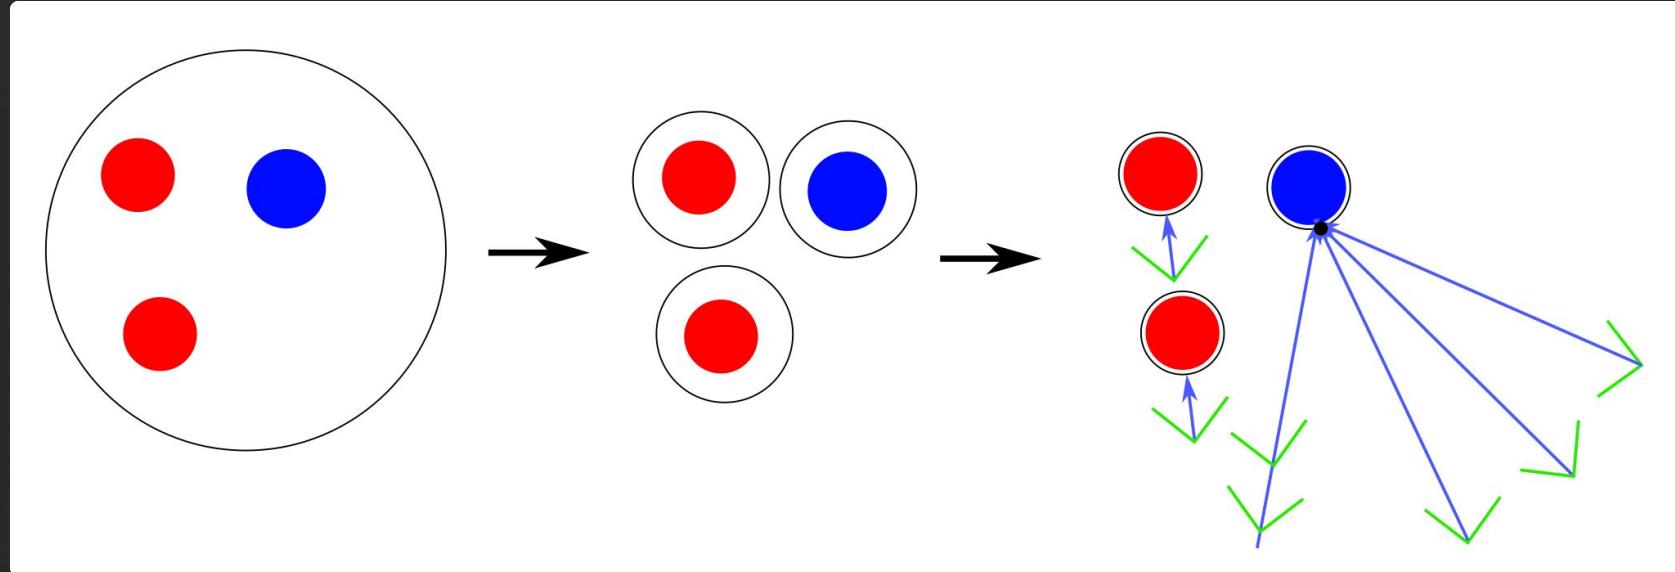


Problems with the Light Field

We fixed redundancy along ray...
But many related rays are far apart in θ, ϕ
(for Lambertian scenes)
→ Hard to learn a regression



Let's Invent...



Single light field

Multiple light fields

Surface light field

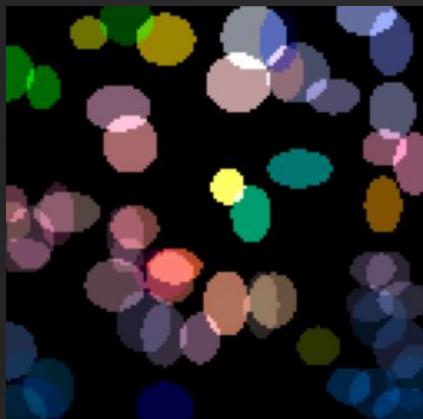
Occlusion problem fixed

Similar rays have similar parameterisations

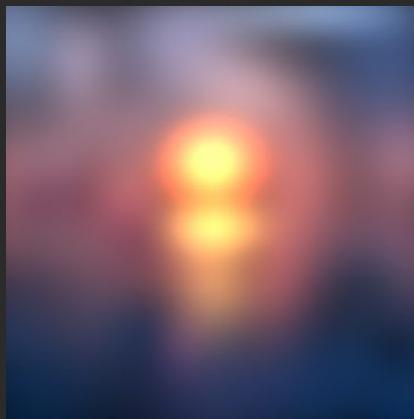
... but how do we pick the surface? i.e. learn the ray parameterisation?

Learning a Light Field Ray Parameterisation

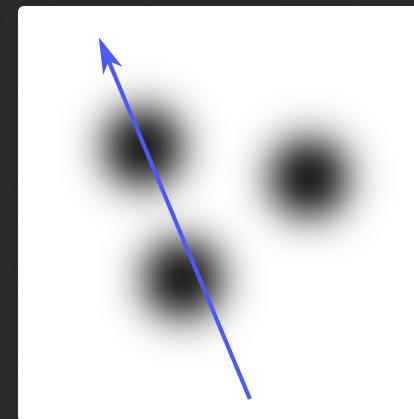
Hard edges?



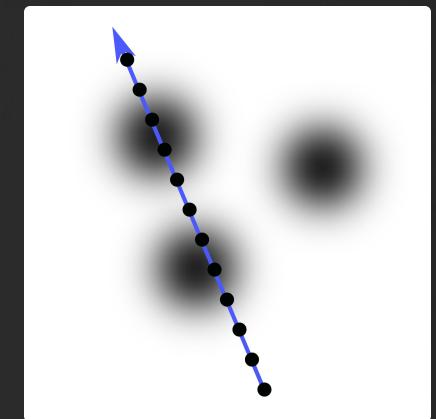
Soft blobs!



Where's the
surface?



Along the
whole ray!



Using our Distributed Light Field Parameterisation

At each point along a ray

Evaluate the local LF

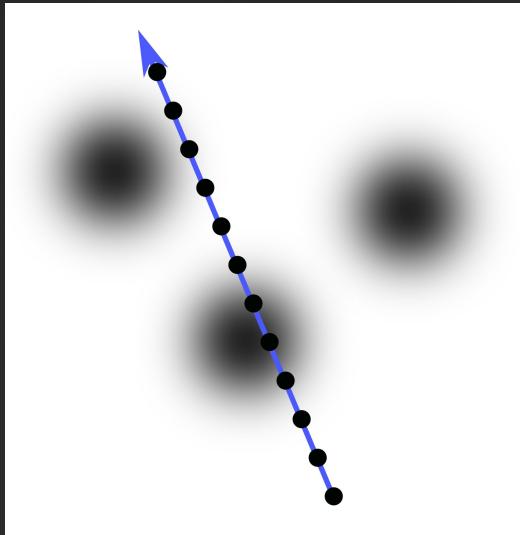
Density $d(x, y)$

Color $c(x, y, \theta)$

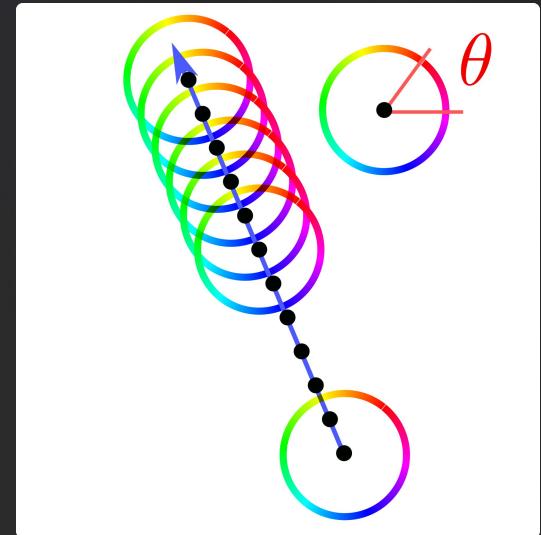
“Density” means how much does
this point contribute visually

A soft surface light field!

We need to account for occlusion...
Easiest to understand approach:



Density



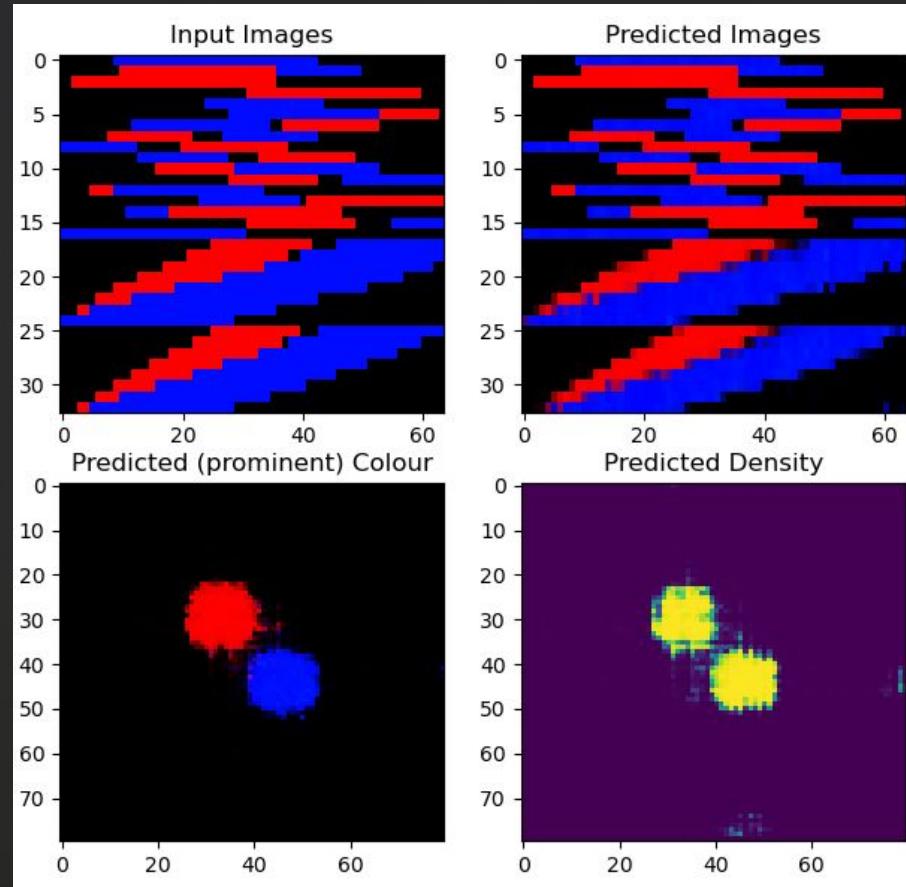
Colour

Along the ray from far to near
Take new colour proportional to density
 $C_i = c \cdot d + C_{i-1} \cdot (1 - d)$

Let's Neural Regress the Soft Surface Light Field

Neurally regress density $d(x, y)$ and color $c(x, y, \theta)$: nvs_nn_nerf.py

Test: Success!



Let's Get Explicit: Gaussian Primitives

Density = \sum Gaussians(x, y)

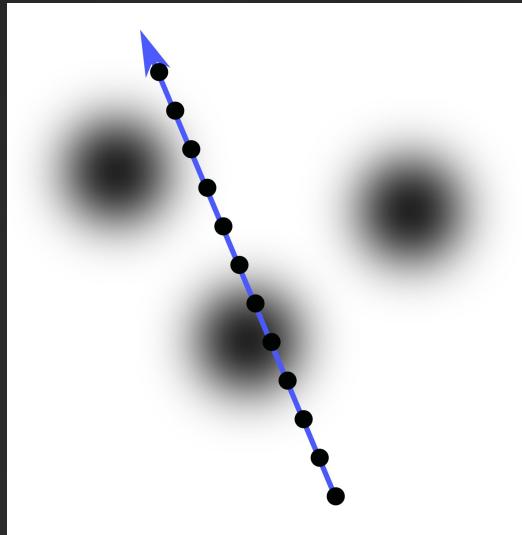
Just like our 2D sunset renders!

Blob has position, size, density

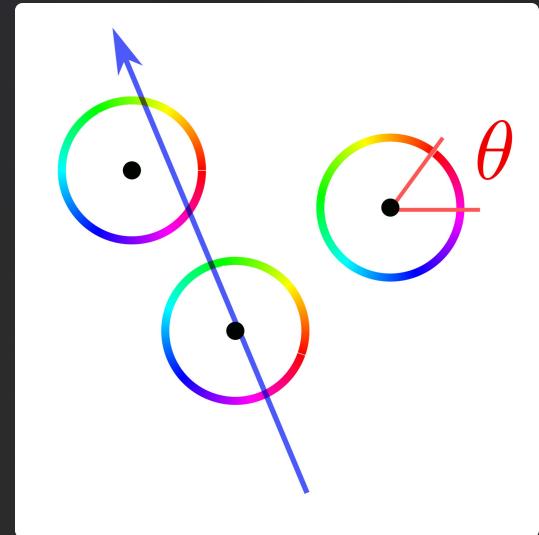
What about color $c(x, y, \theta)$?

We could store a sampled function $c(x, y, \theta)$.

Why not store a single $c(\theta)$ with each Gaussian?



Density



Colour

And compose $c(\theta)$ with circular harmonics?

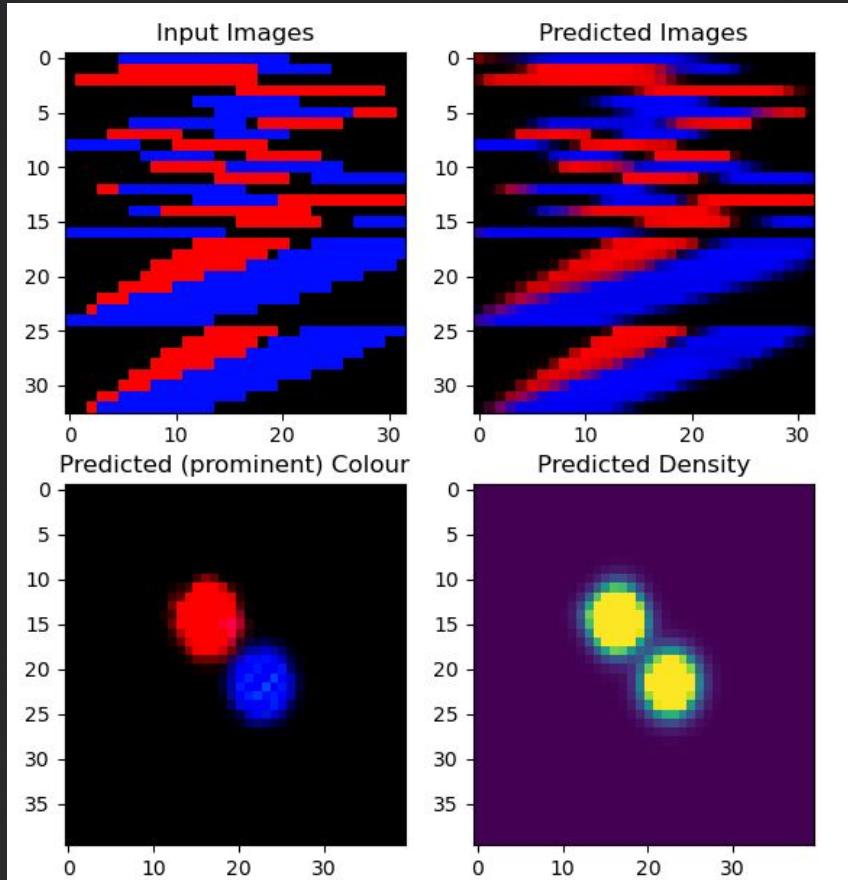
$$c(\theta) = \sum A_n \cos(n \cdot \theta) + B_n \sin(n \cdot \theta)$$

$$n = \dots, -2, -1, 0, 1, 2, \dots$$

Ray Marching on Gaussian Primitives

Raymarch Gaussians $d(x, y)$ with harmonics $c(\theta)$: nvs_gaussian_raymarch.py

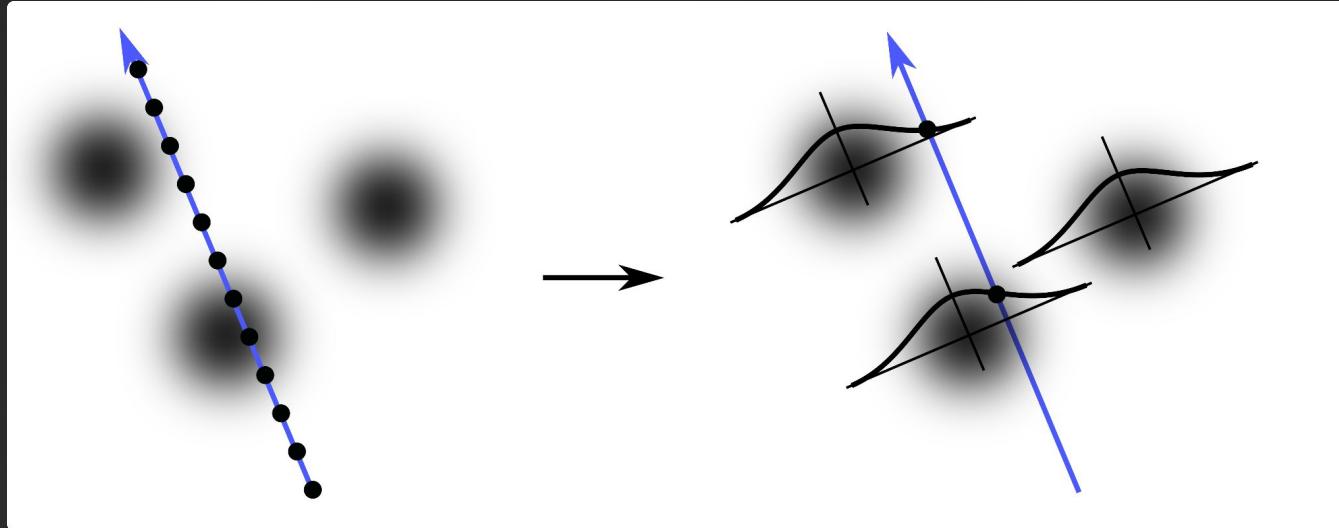
Test: Success! (Sometimes)



To try:

- Anisotropy
- Start without harmonics
- Different init params
- Periodically reduce alpha
- Regularisers

Taking a Shortcut



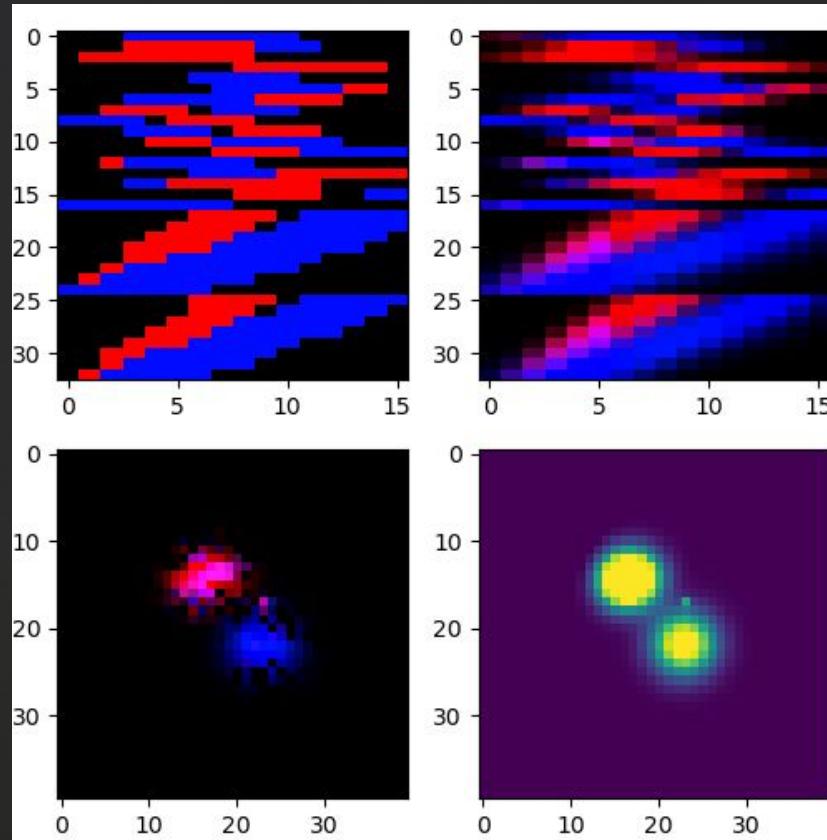
Many evaluations
for every ray
for every blob in 2D

Evaluate each blob once
At densest point
Fewer evaluations
Same blending
Much faster
Approximate

Splatting Gaussian Primitives

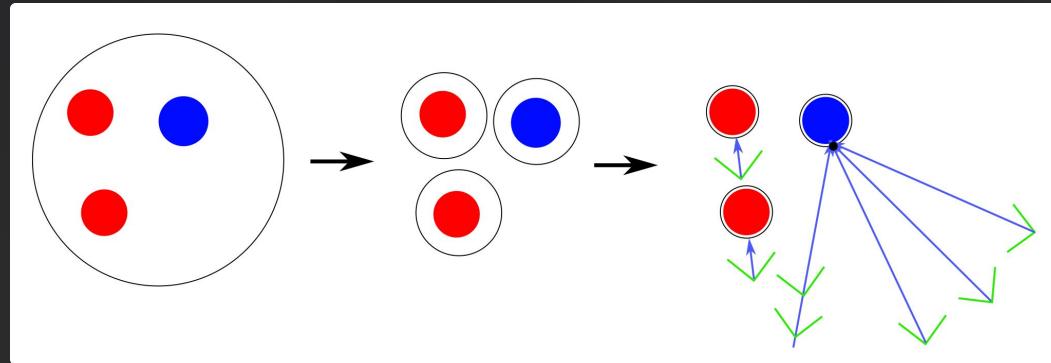
Splat Gaussians for $d(x, y)$ with harmonics $c(\theta)$: nvs_gaussian_splat.py

Test: Success! (*Sometimes*)



Recap and Terminology

Surface light field
Distributed density + colour

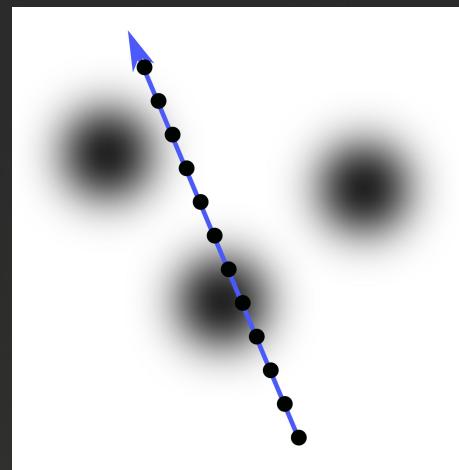


In 3D: “Radiance Field”

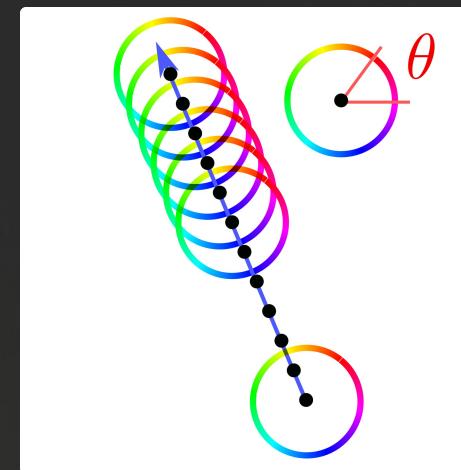
Occlusion handling:

$$C_i = c \cdot d + C_{i-1} \cdot (1 - d)$$

“Alpha blending”



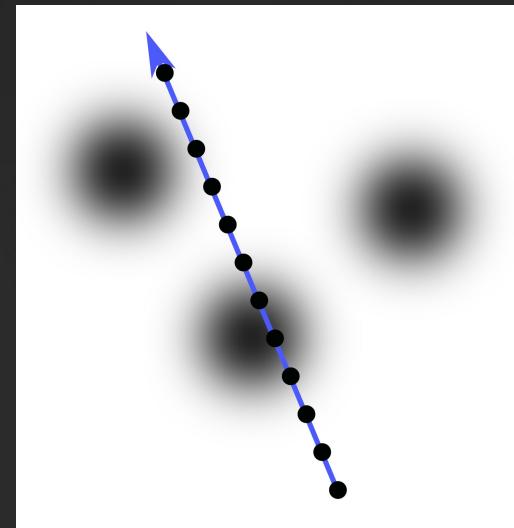
2D: Density $d(x, y)$
3D: Density $d(x, y, z)$



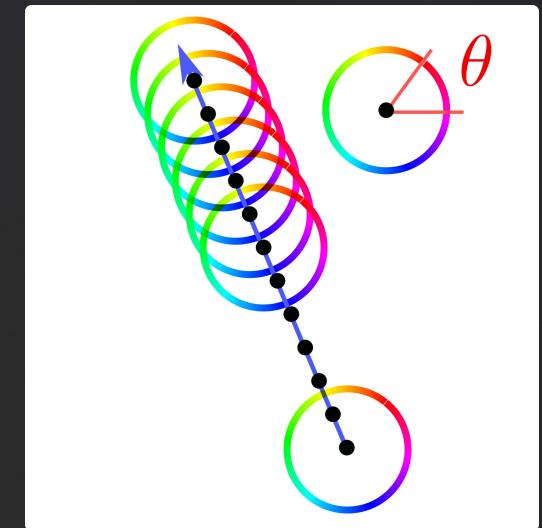
Colour $c(x, y, \theta)$
Colour $c(x, y, z, \theta, \phi)$

Recap and Terminology

Neurally regress density, colour
“Neural Radiance Field”, NeRF



2D: Density $d(x, y)$
3D: Density $d(x, y, z)$



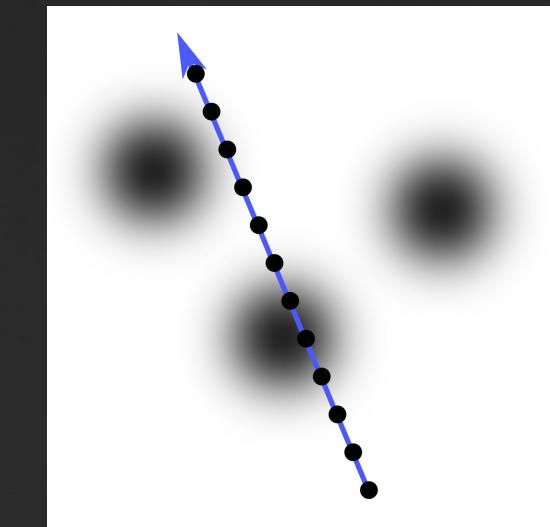
Colour $c(x, y, \theta)$
Colour $c(x, y, z, \theta, \phi)$

Recap and Terminology

Build density with Gaussians

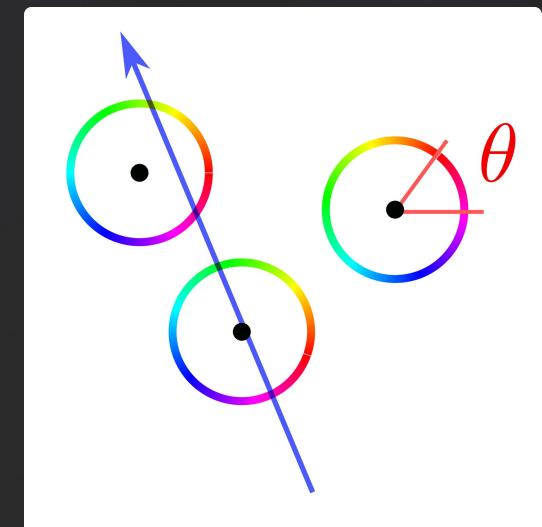
Per-blob colour with harmonics

*"Gaussian Splatting", GS3D
(almost)*



$$2D: d = \sum \text{Gauss}(x, y)$$

$$3D: d = \sum \text{Gauss}(x, y, z)$$



$$c(\theta) = \sum \text{circular harmx}$$

$$c(\theta, \phi) = \sum \text{spherical h.}$$

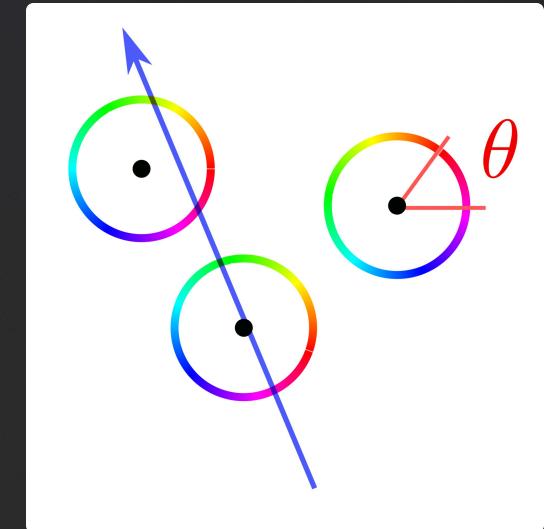
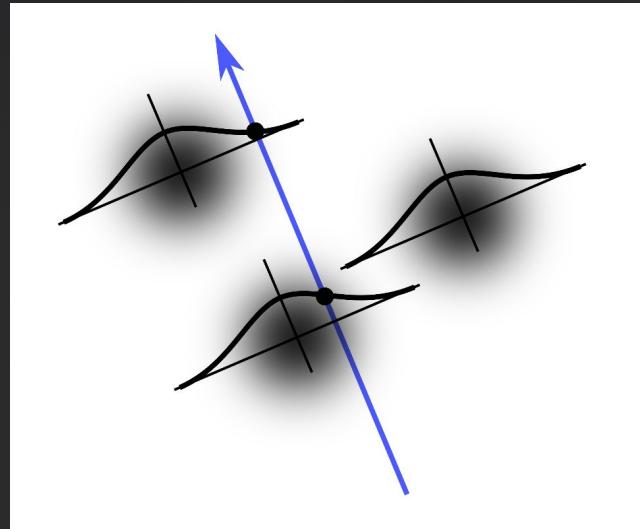
Recap and Terminology

Build density with Gaussians

Per-blob colour with harmonics

Approximate via splatting

Gaussian Splatting, GS3D



2D: Evaluate each blob once

3D: Evaluate each blob once

“Splatting”? “Alpha blending”? Whence these terms?

“(Incomplete) History of Points”, a SIGGRAPH tutorial

Particle systems [Reeves]

Points as a display primitive [Whitted, Levoy]

Oriented particles [Szeliski, Tonnesen]

Particles and implicit surfaces [Witkin, Heckbert]

Surfels [Pfister et al.]

QSplat [Rusinkiewicz, Levoy]

Point Clouds [Linsen, Prautzsch]

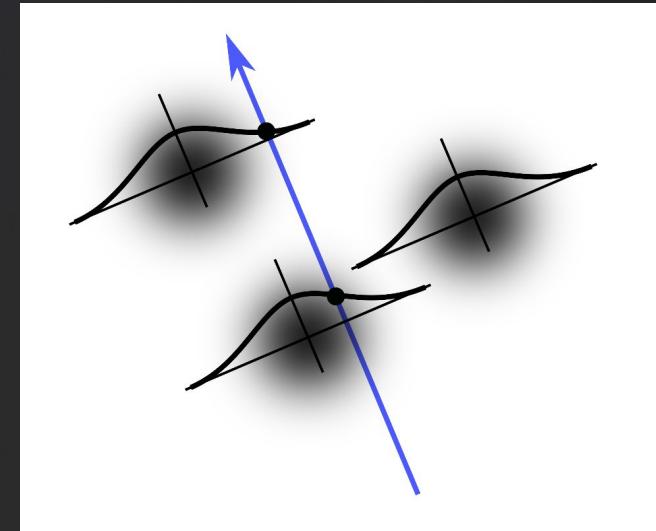
Point set surfaces [Alexa et al.]

Radial basis functions [Carr et al.]

Surface splatting [Zwicker et al.]

Opacity hulls [Matusik et al.]

Sequential Point Trees [Dachsbacher, Stamminger]



$$C_i = c \cdot d + C_{i-1} \cdot (1 - d)$$

GUESS THE YEAR

The Future Lies in the Past!

(Incomplete) History of Points, from a SIGGRAPH tutorial 2004

Particle systems [Reeves 1983]

Points as a display primitive [Whitted, Levoy 1985]

Oriented particles [Szeliski, Tonnesen 1992]

Particles and **implicit surfaces** [Witkin, Heckbert 1994]

Surfels [Pfister et al. 2000]

QSplat [Rusinkiewicz, Levoy 2000]

Point Clouds [Linsen, Prautzsch 2001]

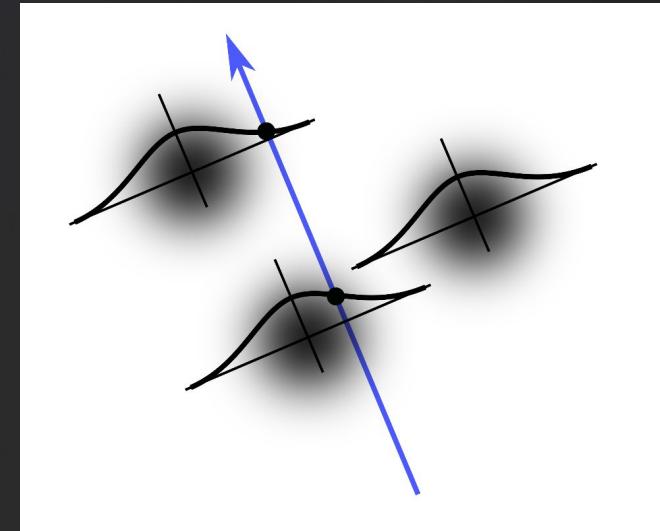
Point set surfaces [Alexa et al. 2001]

Radial basis functions [Carr et al. 2001]

Surface splatting [Zwicker et al. 2001]

Opacity hulls [Matusik et al. 2002]

Sequential Point Trees [Dachsbacher, Stamminger 2003]



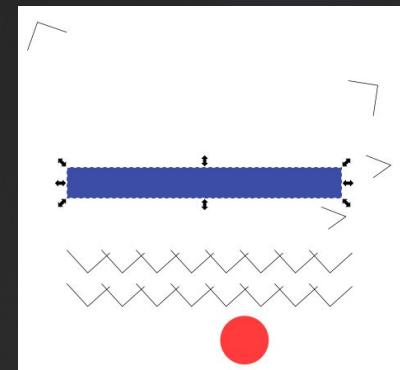
$$C_i = c \cdot d + C_{i-1} \cdot (1 - d)$$

BONUS: A Study in Complex Appearance

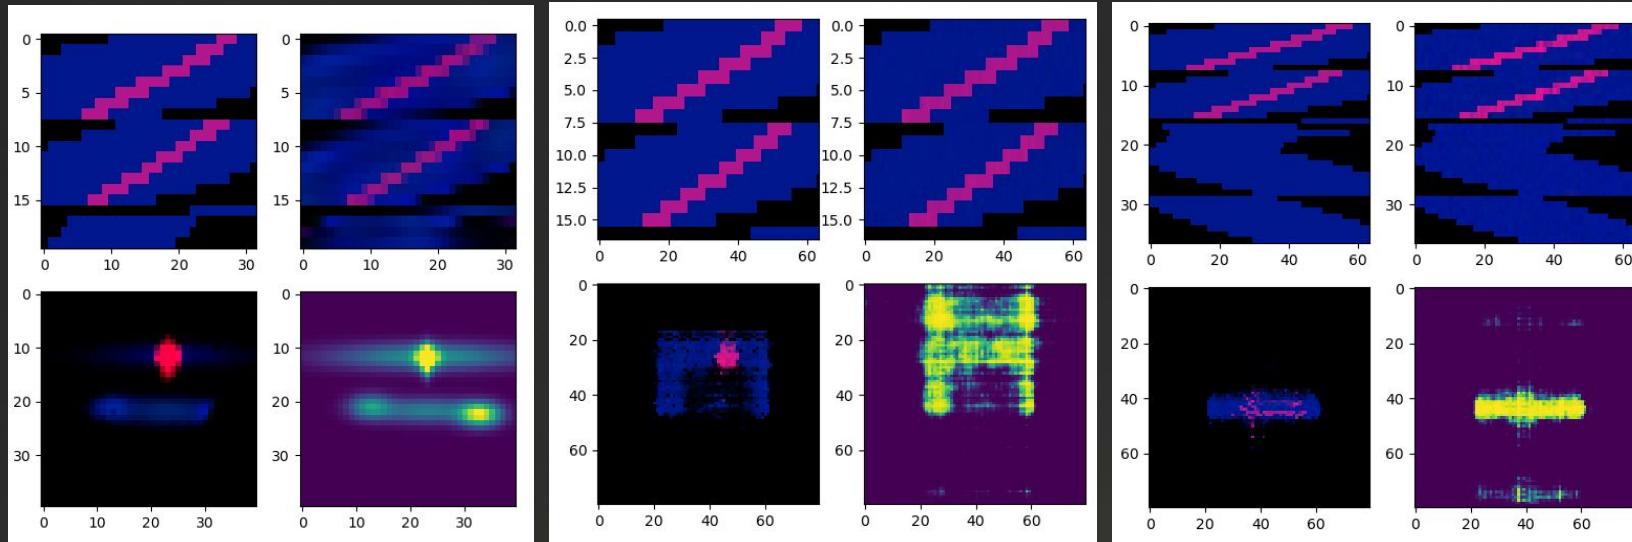
Let's add reflections to our renderer
And fit some radiance fields

scene_refl.svg

Broad range of behaviours:



XML Editor x	
Reflect	0.7
y	132.79042
x	21.235098
height	13.944242
width	125.84338



Radiance Fields are Under-Constrained

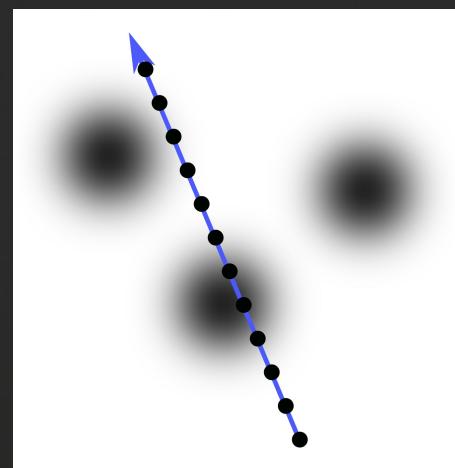
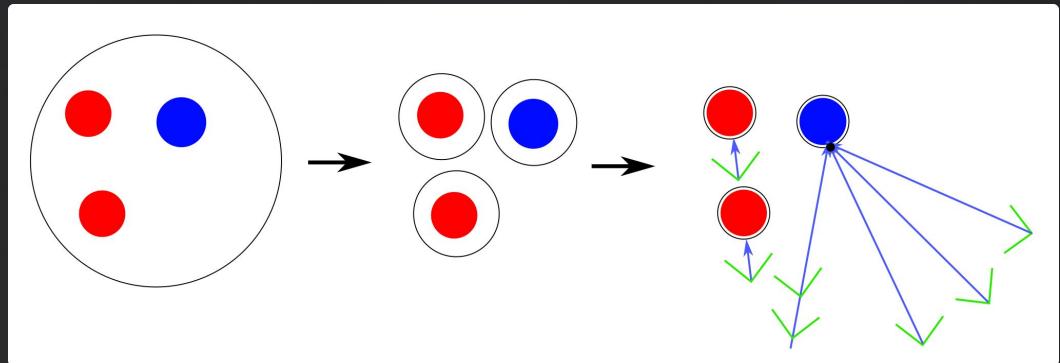
Radiance field is a distributed light field

Parameterisation is a free choice

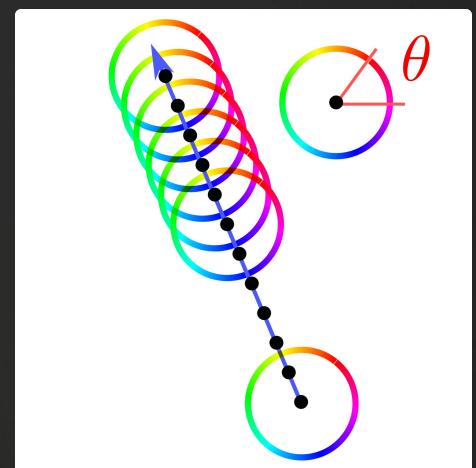
Some easier to train

Learning introduces inductive bias

Many radiance fields have identical appearance



Density



Colour

Checking in 3D

GS3D



NeRF



TODO: explore up-close in-browser GS3D: <https://gsplat.tech/>

Summary: What is a Radiance Field?

Representation of light

Decomposes the 5D plenoptic function into two pieces:

Density in 3D space

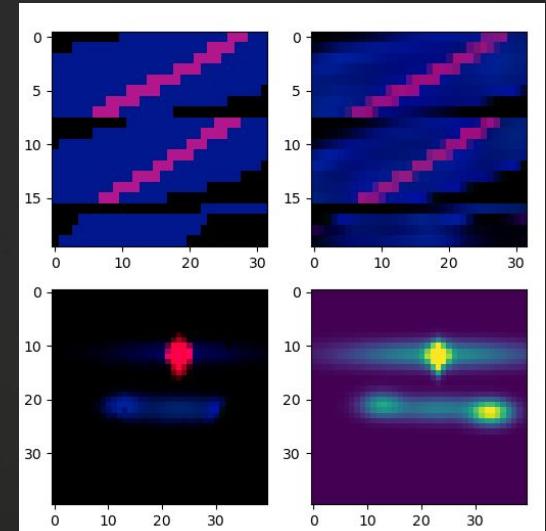
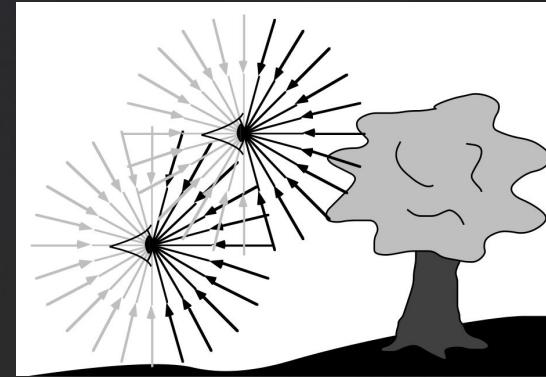
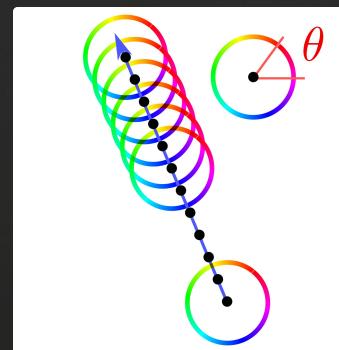
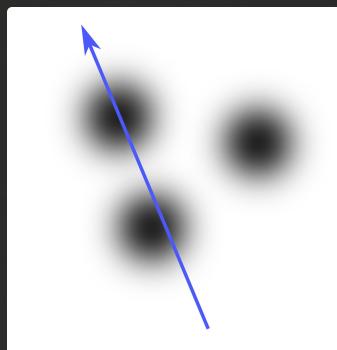
Colour in 2D direction

As a soft, distributed light field suitable for inverse rendering

Natively handles complex appearance in a photo-realistic way!

Density is free, visual parameterisation, not physical density

Many radiance fields have *identical* appearance



Summary: What are NeRF and 3DGS?

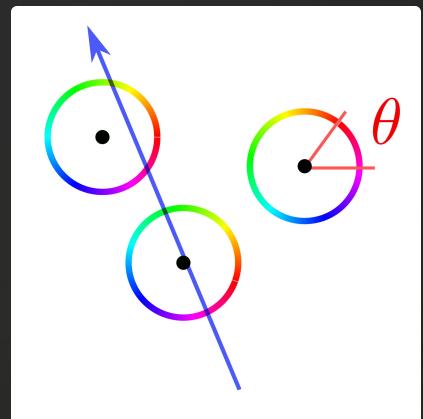
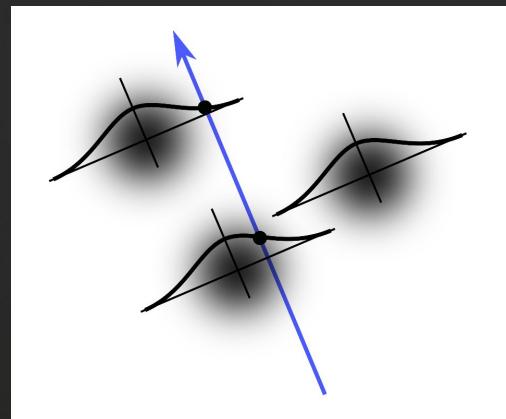
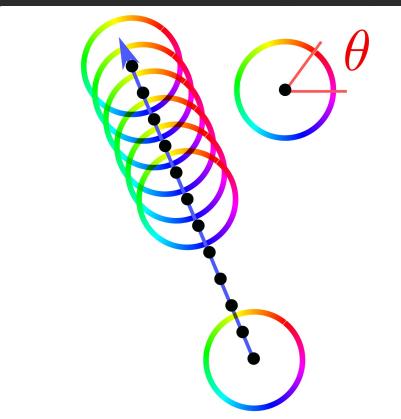
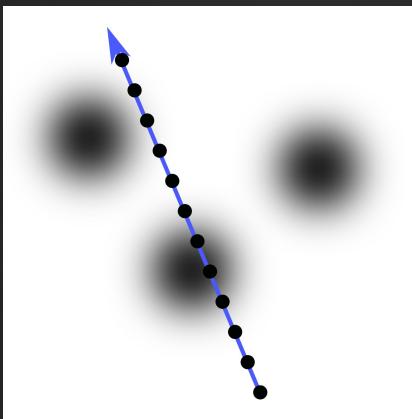
Inverse image-based renderers

Approximate a radiance field to model the world

Via neural regression or by splatting Gaussian primitives

Natively handle complex appearance in a photo-realistic way!

Many NNs and splats yield *similar or identical* radiance fields



Questions and Further Reading

Exciting to summarise many images in a photo-realistic way... but:

How can I analyse what's in the model to understand the scene?

How do I know the camera poses?

Can I build incrementally?

How do I add new images?

How / when do I remove old images?

Should I keep the model and throw away (some?) images?

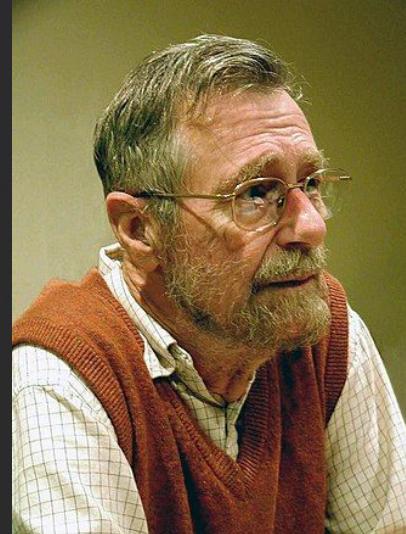
...

Freitas, Davi R., et al. "A Comparative Assessment Of Implicit and Explicit Plenoptic Scene Representations." 2024 IEEE 26th International Workshop on Multimedia Signal Processing (MMSP). IEEE, 2024.

A Parting Thought

“Computer science is no more about computers
than astronomy is about telescopes.”

-- Edsger Dijkstra
(maybe)



Elevates computer science to a legitimate science

Also elevates the role of computers and programs

Your code is a scientific instrument

Small, fast examples build understanding and intuition