

# LEARNING TO ACT

Session 2: Imitation Learning (+RL)

Dana Kulić  
Monash University  
[dana.kulic@monash.edu](mailto:dana.kulic@monash.edu)

RVSS 2026

# Session 2 Overview

- Imitation Learning
  - Inverse Reinforcement Learning
  - Behavioural Cloning
- A brief peek at reinforcement learning

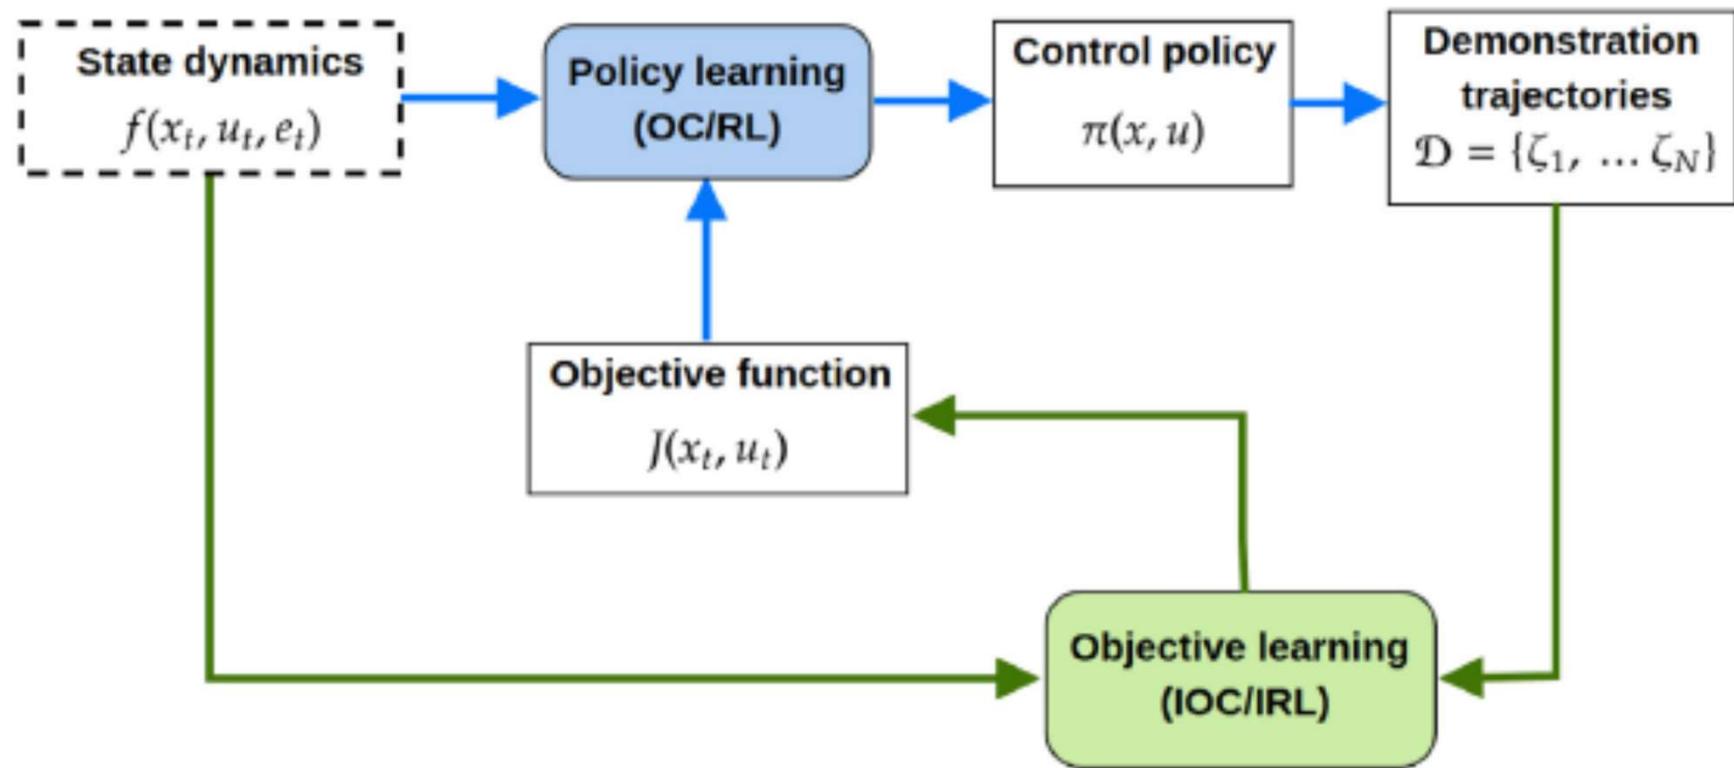
# Review of policy methods

Method	Transition Function	Reward Function	Expert Demos
Optimal Control	Known	Known	No
Reinforcement Learning	Known or Unknown	Known	No
Imitation Learning – IRL	Known or Unknown	Unknown	Yes
Imitation Learning – BC	Known or Unknown	Implicit	Yes

# Imitation Learning

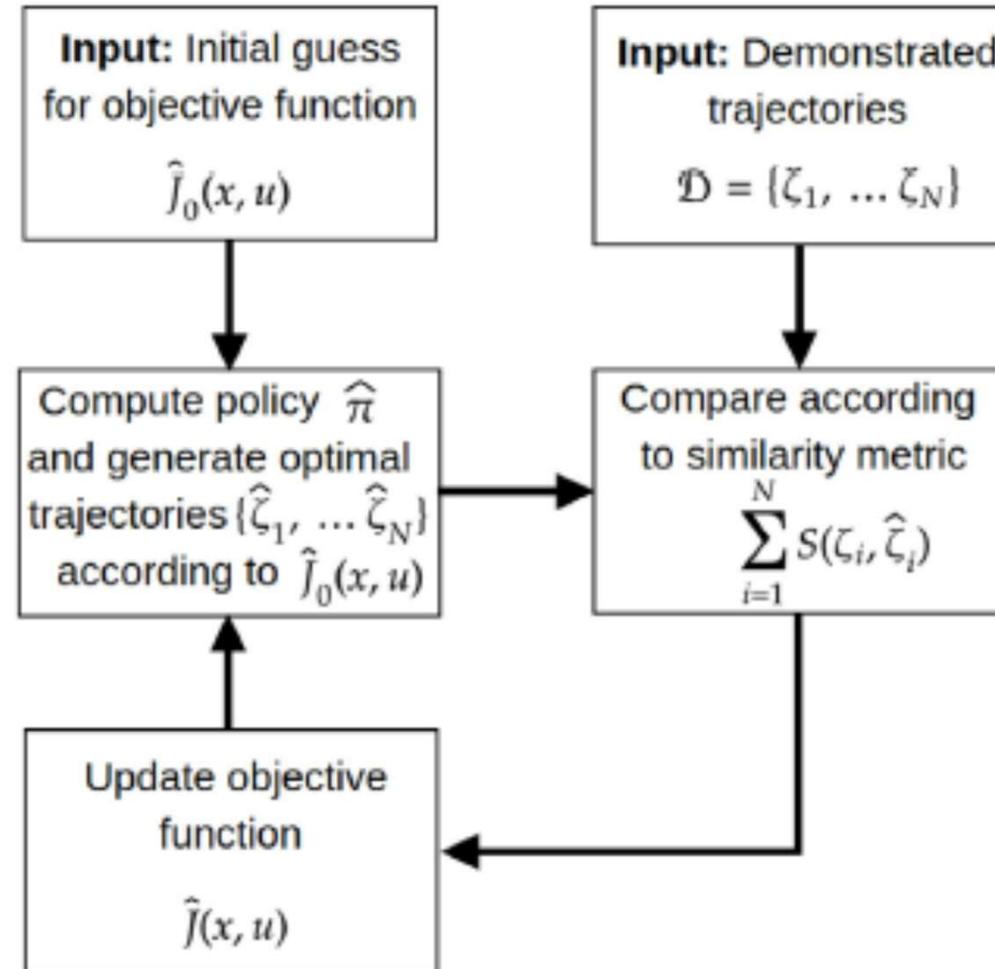
- What if the reward is hard to formulate?
- And I have an expert who can show me the best way to do the task?
- Two approaches:
  - Inverse Reinforcement Learning (IRL): Use the human demonstrations to learn the reward function, then use RL or OC to learn the policy
  - Behaviour Cloning (BC): Learn the policy directly

# Learning the Objective (Cost) Function



J. F. S. Lin, P. Carreno-Medrano, M. Parsapour, M. Sakr, D. Kulić, Objective learning from human demonstrations, Annual Reviews in Control, 2021.

# Bi-level objective learning



J. F. S. Lin, P. Carreno-Medrano, M. Parsapour, M. Sakr, D. Kulić, Objective learning from human demonstrations, Annual Reviews in Control, 2021.

# Behavioural Cloning

- ▶ Collect data from demonstration episodes  $\mathcal{D}(e_{1:N})$
- ▶ Each episode is a sequence of states and actions  
 $e_i = (s_0, a_1, s_1, a_2, \dots, s_T)$
- ▶ Learn a policy  $\phi(s)$  using supervised learning:

$$L = (a_{\mathcal{D}}(s) - \phi(s))^2$$

- ▶ The state  $s$  corresponds to the input data
- ▶ The action  $a$  corresponds to the label
- ▶ Behavioural cloning learns the policy function  $\phi(s)$  to minimise the difference between the estimated action and the observed expert action from each state

# Behavioural Cloning

---

**Algorithm 1** Abstract of behavioral cloning

---

Collect a set of trajectories demonstrated by the expert  $\mathcal{D}$

Select a policy representation  $\pi_\theta$

Select an objective function  $\mathcal{L}$

Optimize  $\mathcal{L}$  w.r.t. the policy parameter  $\theta$  using  $\mathcal{D}$

**return** optimized policy parameters  $\theta$

---

## Key Design Choices:

- What loss function should be used to represent the difference between demonstrated and produced behaviour?
- How should the state be represented?
- How should the policy be represented?
- Output a single action or a trajectory?

# Behavioural Cloning

- What might be some problems we encounter when we apply behavioural cloning?

# Behavioural Cloning – potential problems

- How to collect the expert data?
- What is the right state representation? Does the robot see the same things as the expert does?
- Expert demonstrations may cover only a very small region of the state-space
  - For large state/action spaces, may require a huge data collection effort
- What should the robot do when it encounters a situation that wasn't seen in the dataset?
- How to handle variations in strategy?

# Behavioural Cloning – potential problems

- **How to collect the expert data?**
- What is the right state representation? Does the robot see the same things as the expert does?
- Expert demonstrations may cover only a very small region of the state-space
  - For large state/action spaces, may require a huge data collection effort
- What should the robot do when it encounters a situation that wasn't seen in the dataset?
- How to handle variations in strategy?

# How to collect expert data?



From human demonstrations directly



Via teleoperation

# How to collect demonstration data?



# Properties of demonstration data

- Do I get the state only or both the state and action?
- How valid is the assumption that the expert data is optimal?
  - Intuitive interface
  - Sensor noise
  - Demonstrator expertise
- How much of the state-space/task variation is covered?
- Are there variations in strategy?

Are there other ways to collect expert input?

# Are there other ways to collect expert input?

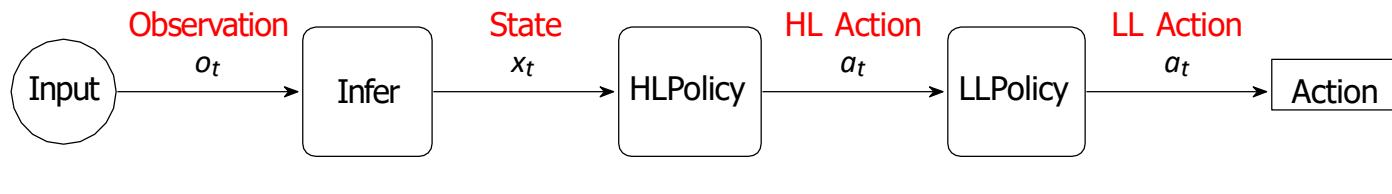
- Preference Learning:
  - Use preferences to learn the human reward function - Reinforcement Learning from Human Feedback (RLHF, [Christiano et al. 2017](#))
  - Use preferences to learn the policy directly – Direct Preference Optimisation (DPO – [Rafailov et al. 2023](#))
  - Make the process more efficient by encoding prior on likely human biases – Kahneman-Tversky Optimisation ([KTO – Ethayarajh et al. 2024](#))
- Unifying multiple human signals for learning ([Jeon et al., 2020](#))
- Use (simulated) optimisation solutions as demonstrations

# Behavioural Cloning – potential problems

- How to collect the expert data?
- **What is the right state representation? Does the robot see the same things as the expert does?**
- Expert demonstrations may cover only a very small region of the state-space
  - For large state/action spaces, may require a huge data collection effort
- What should the robot do when it encounters a situation that wasn't seen in the dataset?
- How to handle variations in strategy?

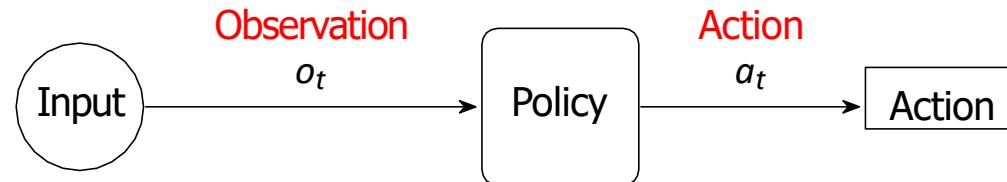
# Choices for Problem formulation

Modular



18

End-to-end



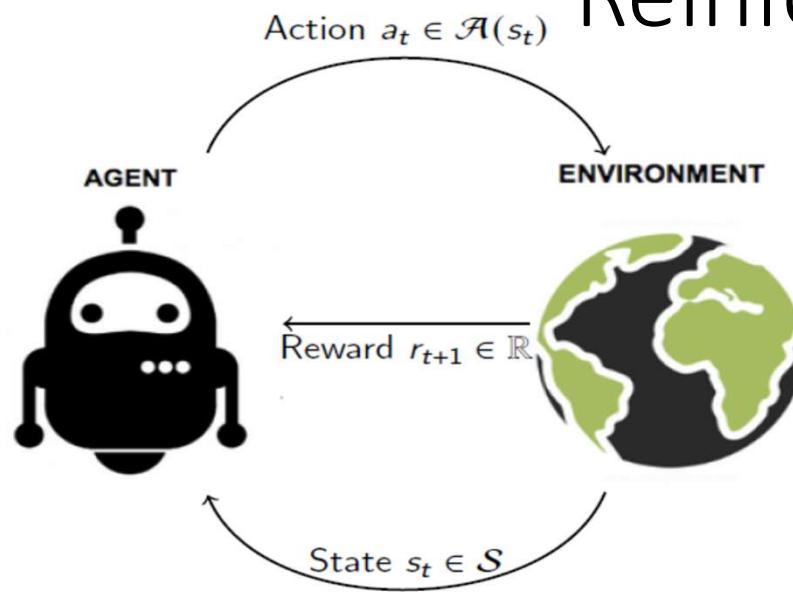
# Behavioural Cloning – potential problems

- How to collect the expert data?
- What is the right state representation? Does the robot see the same things as the expert does?
- Expert demonstrations may cover only a very small region of the state-space
  - For large state/action spaces, may require a huge data collection effort
- What should the robot do when it encounters a situation that wasn't seen in the dataset?
- How to handle variations in strategy?

# Let's try it out!

- Please open `IntroBC.ipynb` from the `Reinforcement_Learning` folder in the school repo

# Reinforcement Learning



The agent and the environment interact at discrete time steps  $t = \{1, 2, \dots\}$ . At each step  $t$ ,

the agent:

- ▶ Observes state  $s_t$
- ▶ Executes action  $a_t$
- ▶ Receives scalar reward  $r_{t+1}$

the environment:

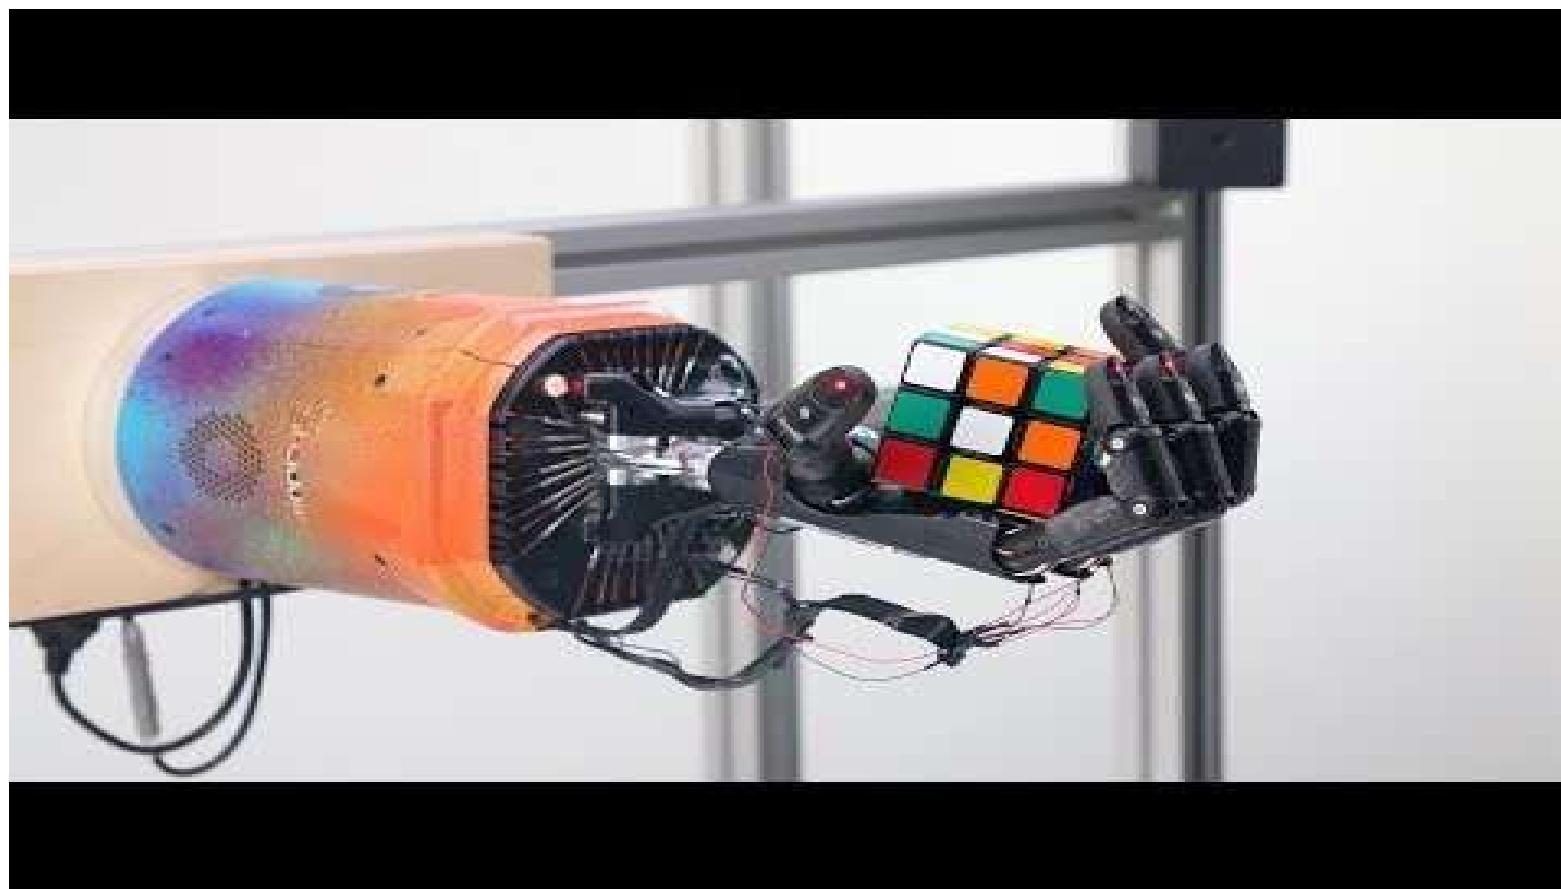
- ▶ Receives action  $a_t$
- ▶ Emits state  $s_{t+1}$
- ▶ Emits scalar reward  $r_{t+1}$

Image Courtesy of [Lilian Weng](#)

# RL Examples



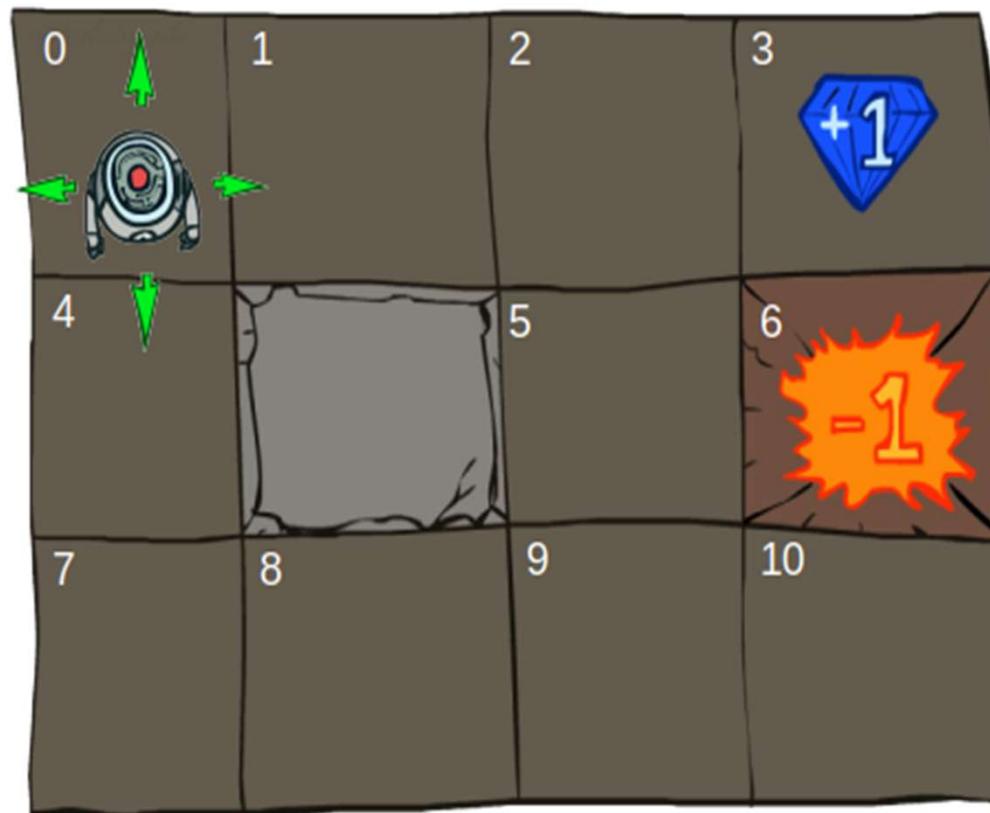
# RL Examples



# RL Examples



The problem with maximising reward from experience



# Value Functions

- A value function defines the amount of reward an agent can expect to accumulate over the future under a particular policy

The *state-value* function  $v_\pi(s)$  is the expected return when starting in state  $s$  and following  $\pi$  thereafter,

$$v_\pi(s) = \mathbb{E}_\pi \left[ \sum_{k=0}^{\infty} \gamma^k r_{t+k+1} | s_t = s \right], \forall s \in \mathcal{S}$$

# Value Functions and Bellman Equations

- The state-value function can be decomposed into the immediate reward plus the discounted value of the successor state

$$v_\pi(s) = \mathbb{E}_\pi \left[ \sum_{k=0}^{\infty} \gamma^k r_{t+k+1} | s_t = s \right], \forall s \in \mathcal{S}$$

$$\begin{aligned} v_\pi(s) &= \mathbb{E}_\pi \left[ r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots | s_t = s \right] \\ &= \mathbb{E}_\pi \left[ r_{t+1} + \gamma (r_{t+2} + \gamma r_{t+3} + \dots) | s_t = s \right] \\ &= \mathbb{E}_\pi \left[ \underbrace{r_{t+1}}_{\text{Immediate reward}} + \overbrace{\gamma v_\pi(s_{t+1})}^{\text{Discounted value}} | s_t = s \right] \end{aligned}$$

# Optimal Value and Policy

The optimal value function produces the maximum return:

$$V_*(s) = \max_{\pi} V_{\pi}(s),$$

The optimal policy achieves optimal value functions:

$$\pi_* = \arg \max_{\pi} V_{\pi}(s);$$

# How does this help us to learn the best policy?

- Start with some policy guess
- Estimate how good that policy is (by estimating its value function)
- Improve the policy
- Iterate -> Policy Iteration

$$\pi_0 \xrightarrow{E} v_{\pi_0} \xrightarrow{I} \pi_1 \xrightarrow{E} v_{\pi_1} \xrightarrow{I} \pi_2 \xrightarrow{E} \dots \xrightarrow{I} \pi^* \xrightarrow{E} v^*$$

# Algorithm for Policy Evaluation

---

**Algorithm 1:** Iterative Policy Evaluation for estimating  $v \approx v_\pi$

---

**Input :** MDP tuple  $\langle \mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \gamma \rangle$ , policy  $\pi$ , threshold  $\theta > 0$

**Output:**  $v_\pi(s)$

Initialize  $v(s) = 0 \forall s \in \mathcal{S}$

**repeat**

$\Delta \leftarrow 0$

**foreach**  $s \in \mathcal{S}$  **do**

$V \leftarrow v(s)$

$v(s) \leftarrow \sum_{a \in \mathcal{A}} \pi(a|s) \left( \mathcal{R}(s, a) + \gamma \sum_{s' \in \mathcal{S}} \mathcal{T}(s, a, s') v(s') \right)$

$\Delta \leftarrow \max(\Delta, |V - v(s)|)$

**until**  $\Delta < \theta$

---

# Algorithm for Policy Improvement

---

**Algorithm 2:** Policy Iteration for estimating  $\pi \approx \pi^*$

---

**Input :** MDP tuple  $\langle \mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \gamma \rangle$

**Output:**  $\pi \approx \pi^*$

Initialize  $\pi(s) \forall s \in \mathcal{S}$  to a random action  $a \in \mathcal{A}$ , arbitrarily

**repeat**

$\pi' \leftarrow \pi$

    Compute  $v_\pi(s)$  for all states using *policy evaluation*

**foreach**  $s \in \mathcal{S}$  **do**

$\pi(s) \leftarrow \arg \max_{a \in \mathcal{A}} \mathcal{R}(s, a) + \gamma \sum_{s' \in \mathcal{S}} \mathcal{T}(s, a, s') v_\pi(s')$

**until**  $\pi(s) == \pi'(s) \forall s \in \mathcal{S}$

---

# Let's try it out!

- Please open `IntroRL.ipynb` from the `Reinforcement_Learning` folder in the school repo

# What if we don't have the model? – Temporal Difference Learning

**Goal:** Given a policy  $\pi$ , learn  $\hat{v}_\pi(s)$  from experience episodes

$$\{s_0, a_0, r_1, \dots, s_T\} \sim \pi$$

**Recall:** State-value Bellman Equation

$$v(s_t) = \mathbb{E}_\pi \left[ r_{t+1} + \gamma v_\pi(s_{t+1}) | s_t = s \right]$$

**Approximation:** At each time step  $t$  use *observed immediate* reward  $r_{t+1}$  and the estimated return  $\hat{v}(s_{t+1})$  to update  $\hat{v}(s_t)$

$$\hat{v}(s_t) \leftarrow \hat{v}(s_t) + \alpha \underbrace{[r_{t+1} + \gamma \hat{v}(s_{t+1}) - \hat{v}(s_t)]}_{\text{TD target}}$$

where  $\alpha$  is a step-size parameter.

We update our sample estimate  $\hat{v}(s_t)$  in the direction of the TD error.

# Q-learning

---

**Algorithm 2:** Q-learning algorithm

---

**Input :** Step-size  $\alpha \in (0, 1]$ , small  $\epsilon > 0$

**Output:**  $\hat{q}^*(s, a)$

Initialize  $\hat{q}(s, a)$  arbitrarily  $\forall s \in \mathcal{S}, a \in \mathcal{A}$ ,  $q(\text{terminal state}, \cdot) = 0$

**Loop** for each episode

    Initialize  $s$

**repeat**

        Choose action  $a$  from  $s$  using  $\epsilon$ -greedy policy derived from  
 $\hat{q}(s, a)$

        Take action  $a$ , observe  $r, s'$

$\hat{q}(s, a) \leftarrow \hat{q}(s, a) + \alpha [r + \gamma \max_{a'} \hat{q}(s', a') - \hat{q}(s, a)]$

$s \leftarrow s'$

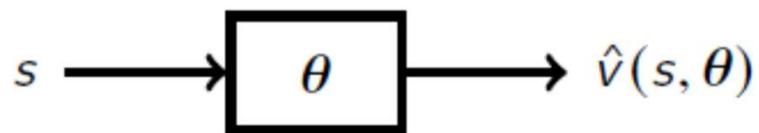
**until**  $s$  is terminal

---

# Let's try it out!

- Please open TD-RL.ipynb from the Reinforcement\_Learning folder in the school repo

# Function Approximation



$$\hat{v}(s, \theta) \approx v_\pi(s)$$



$$\hat{q}(s, a, \theta) \approx q_\pi(s, a)$$

# Function Approximation

- If we knew the true action-value function, this would be a standard supervised learning problem, we could find the best approximation by minimising:

$$J(\theta) = \mathbb{E}_\pi [(q_\pi(s, a) - \hat{q}_\pi(s, a, \theta))^2]$$

- But RL only gives access to rewards. Use Bellman equation

$$\Delta\theta = \alpha \underbrace{[r_{t+1} + \gamma \hat{q}_\pi(s_{t+1}, a_{t+1}, \theta) - \hat{q}_\pi(s_t, a_t, \theta)]}_{\text{Target}} \nabla_\theta \hat{q}_\pi(s_t, a_t, \theta)$$

# DQN Algorithm

---

**Algorithm 1** Deep Q-learning with Experience Replay

---

Initialize replay memory  $\mathcal{D}$  to capacity  $N$

Implemented in main loop

Initialize action-value function  $Q$  with random weights

**for** episode = 1,  $M$  **do**

    Initialise sequence  $s_1 = \{x_1\}$  and preprocessed sequenced  $\phi_1 = \phi(s_1)$

**for**  $t = 1, T$  **do**

        With probability  $\epsilon$  select a random action  $a_t$

        otherwise select  $a_t = \max_a Q^*(\phi(s_t), a; \theta)$

        Execute action  $a_t$  in emulator and observe reward  $r_t$  and image  $x_{t+1}$

        Set  $s_{t+1} = s_t, a_t, x_{t+1}$  and preprocess  $\phi_{t+1} = \phi(s_{t+1})$

        Store transition  $(\phi_t, a_t, r_t, \phi_{t+1})$  in  $\mathcal{D}$

        Sample random minibatch of transitions  $(\phi_j, a_j, r_j, \phi_{j+1})$  from  $\mathcal{D}$

        Set  $y_j = \begin{cases} r_j & \text{for terminal } \phi_{j+1} \\ r_j + \gamma \max_{a'} Q(\phi_{j+1}, a'; \theta) & \text{for non-terminal } \phi_{j+1} \end{cases}$

        Perform a gradient descent step on  $(y_j - Q(\phi_j, a_j; \theta))^2$  according to equation 3

**end for**

**end for**

Implemented in optimize(.) method of  
the Agent class

# Let's try it out!

- Please open DeepRL\_BasicDQN.ipynb from the Reinforcement\_Learning folder in the school repo
- Further exploration: check out DeepRL\_TargetDQN.ipynb

# Types of RL Algorithms

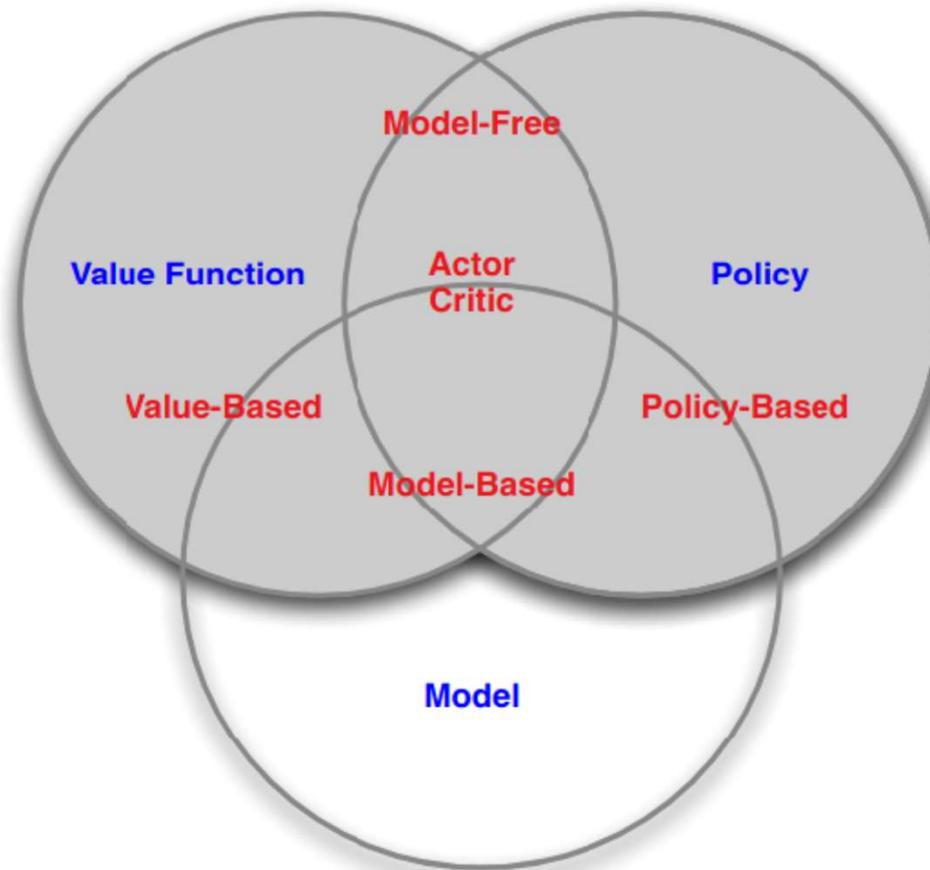


Image Courtesy of [David Silver](#)

# Function approximation for value/policy learning

- Collect data from executions in the replay buffer  $e_t = (S_t, A_t, R_t, S_{t+1})$
- Sample a minibatch of data from the replay buffer
- Update Value function with Bellman loss

$$\mathcal{L}(\theta) = \mathbb{E}_{(s,a,r,s') \sim U(D)} \left[ (r + \gamma \max_{a'} Q(s', a'; \theta^-) - Q(s, a; \theta))^2 \right]$$

- Update policy parameters to maximise the value

$$\mathcal{J}(\theta) = \sum_{s \in \mathcal{S}} d_{\pi_\theta}(s) V_{\pi_\theta}(s) = \sum_{s \in \mathcal{S}} \left( d_{\pi_\theta}(s) \sum_{a \in \mathcal{A}} \pi(a|s, \theta) Q_\pi(s, a) \right)$$

# Challenges of RL

- Exploitation vs. exploration
- Reward Formulation
- “The Deadly Triad”: bootstrapping+off-policy+approximation
- Sim2real gap



# Summary of This Session

- Imitation Learning:
  - Don't need to know the reward
  - Must collect "expert" demonstrations
  - Many design choices about problem formulation
- Reinforcement Learning:
  - Must specify the reward
  - From reward to value
  - From value to policy